

Paging

Abhilash Jindal

Overview


- More flexible address translation with paging (OSTEP Ch 18-20)
 - Paging hardware
- Demand paging: swapping pages to disk when memory becomes full (OSTEP Ch 21-22)
 - Swapping mechanisms
 - Page replacement algorithms
- Paging in action (xv6 book Ch 2, OSTEP Ch 23)
 - Paging on xv6
 - Fork with Copy-on-write, Guard pages

Paging Hardware


OSTEP Ch 18-20

Intel SDM Volume 3A Ch 4


Memory isolation and address space




Memory isolation and address space




Memory isolation and address space




Memory isolation and address space



Segmentation




Segmentation




- Mapping large address spaces

Segmentation




- Mapping large address spaces
- Place each segment independently to not map free space

Segmentation




- Mapping large address spaces
- Place each segment independently to not map free space



Growing and shrinking address space

- Segments need to be contiguous in memory
- Growing might not succeed if there are other segments next to heap segment



Limitations of segmentation




Figure 3-4. Multi-Segment Model

Limitations of segmentation

Limited flexibility:




Figure 3-4. Multi-Segment Model

Limitations of segmentation

Limited flexibility:

- To support large address spaces, burden falls on programmer/compiler to manage multiple segments




Figure 3-4. Multi-Segment Model

Limitations of segmentation

Limited flexibility:

- To support large address spaces, burden falls on programmer/compiler to manage multiple segments
- Only an entire segment can be shared. Example: cannot share some part of CS (e.g., both processes use same library)




Figure 3-4. Multi-Segment Model

Limitations of segmentation

Limited flexibility:

- To support large address spaces, burden falls on programmer/compiler to manage multiple segments
- Only an entire segment can be shared. Example: cannot share some part of CS (e.g., both processes use same library)

Different sized segments, segments need to be contiguous in physical memory




Figure 3-4. Multi-Segment Model

Limitations of segmentation

Limited flexibility:

- To support large address spaces, burden falls on programmer/compiler to manage multiple segments
- Only an entire segment can be shared. Example: cannot share some part of CS (e.g., both processes use same library)

Different sized segments, segments need to be contiguous in physical memory

- complicates physical memory allocator




Figure 3-4. Multi-Segment Model

Limitations of segmentation

Limited flexibility:

- To support large address spaces, burden falls on programmer/compiler to manage multiple segments
- Only an entire segment can be shared. Example: cannot share some part of CS (e.g., both processes use same library)

Different sized segments, segments need to be contiguous in physical memory

- complicates physical memory allocator
- lead to external fragmentation




Figure 3-4. Multi-Segment Model

Limitations of segmentation

Limited flexibility:

- To support large address spaces, burden falls on programmer/compiler to manage multiple segments
- Only an entire segment can be shared. Example: cannot share some part of CS (e.g., both processes use same library)

Different sized segments, segments need to be contiguous in physical memory

- complicates physical memory allocator
- lead to external fragmentation
- growing/shrinking segments is awkward




Figure 3-4. Multi-Segment Model

Paging





Figure 18.3: The Address Translation Process



Notebook analogy

Segmentation



OS

0	1	2	3	4	5
xv6	is	an	OS	for	x86



DB

0	1	2	3
Write	an	SQL	query




Notebook

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
OS:1 DB:7	xv6	is	an	OS	for	x86	Write	an	SQL	query				

Reading S

- Read second letter from 3rd page



- Read second letter from 4th page

Notebook analogy


Paging



0	1	2	3	4	5
xv6	is	an	OS	for	x86




0	1	2	3
Write	an	SQL	query



0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
OS:1 DB:5	0:3,1:4, 2:2,3:8, 4:6,5:10	an	xv6	is	0:9,1:2, 2:7,3:12	for	SQL	OS	Write	x86		query		

Reading S

- Read second letter from 3rd page



- Read second letter from 8th page

Paging

Segmentation	Paging
Large address spaces need multi segment model. Burden falls on programmer/compiler to manage multiple segments	Transparently support large address spaces. Programmer/compiler work with a flat virtual address space

Paging

Segmentation	Paging
Large address spaces need multi segment model. Burden falls on programmer/compiler to manage multiple segments	Transparently support large address spaces. Programmer/compiler work with a flat virtual address space
Different sized segments lead to external fragmentation	Fixed sized pages (4KB, 4MB)

Paging

Segmentation	Paging
Large address spaces need multi segment model. Burden falls on programmer/compiler to manage multiple segments	Transparently support large address spaces. Programmer/compiler work with a flat virtual address space
Different sized segments lead to external fragmentation	Fixed sized pages (4KB, 4MB)
Different sized segments complicate physical memory allocator	

Paging

Segmentation	Paging
Large address spaces need multi segment model. Burden falls on programmer/compiler to manage multiple segments	Transparently support large address spaces. Programmer/compiler work with a flat virtual address space
Different sized segments lead to external fragmentation	Fixed sized pages (4KB, 4MB)
Different sized segments complicate physical memory allocator	
Full segment needs to be contiguous in physical memory	Page is contiguous. Neighbouring addresses (in different pages) may not be contiguous

Paging

Segmentation	Paging
Large address spaces need multi segment model. Burden falls on programmer/compiler to manage multiple segments	Transparently support large address spaces. Programmer/compiler work with a flat virtual address space
Different sized segments lead to external fragmentation	Fixed sized pages (4KB, 4MB)
Different sized segments complicate physical memory allocator	
Full segment needs to be contiguous in physical memory	Page is contiguous. Neighbouring addresses (in different pages) may not be contiguous
Growing/shrinking segments is awkward	Simple: allocate another page, free a page

Paging


Segmentation	Paging
Large address spaces need multi segment model. Burden falls on programmer/compiler to manage multiple segments	Transparently support large address spaces. Programmer/compiler work with a flat virtual address space
Different sized segments lead to external fragmentation	Fixed sized pages (4KB, 4MB)
Different sized segments complicate physical memory allocator	
Full segment needs to be contiguous in physical memory	Page is contiguous. Neighbouring addresses (in different pages) may not be contiguous
Growing/shrinking segments is awkward	Simple: allocate another page, free a page
Address translation hardware has an adder ($va + base$) and a comparator ($va < limit$)	Address translation hardware is much more complicated

How to do page-based address translation?

Page table: Maintain a lookup table for each process

Virtual page number	Physical page number
7	3
6	x
5	x
4	x
3	x
2	4
1	6
0	0

Virtual page number	Physical page number
7	5
6	x
5	x
4	x
3	7
2	2
1	1
0	6




How to do page-based address translation? (2)

Page table bits

- Present bit: valid mapping
- Permission bits: read only, writeable, executable

Virtual page number	Physical page number	Present	Permissions
7	3	Y	rw
6	x	N	
5	x	N	
4	x	N	
3	x	N	
2	4	Y	rw
1	6	Y	rx
0	0	Y	rx



Page table size

- Virtual addresses: 2^{32}

Virtual page number	Physical page number
7	3
6	x
5	x
4	x
3	x
2	4
1	6
0	0

Page table size

- Virtual addresses: 2^{32}
- Size of page = (4KB) = 2^{12}

Virtual page number	Physical page number
7	3
6	x
5	x
4	x
3	x
2	4
1	6
0	0

Page table size

- Virtual addresses: 2^{32}
- Size of page = (4KB) = 2^{12}
- Number of page table entries = $2^{32} / 2^{12} = 2^{20}$

Virtual page number	Physical page number
7	3
6	x
5	x
4	x
3	x
2	4
1	6
0	0

Page table size

- Virtual addresses: 2^{32}
- Size of page = (4KB) = 2^{12}
- Number of page table entries = $2^{32} / 2^{12} = 2^{20}$
- Number of pages in a 4GB DRAM = $2^{32} / 2^{12} = 2^{20}$

Virtual page number	Physical page number
7	3
6	x
5	x
4	x
3	x
2	4
1	6
0	0

Page table size

- Virtual addresses: 2^{32}
- Size of page = (4KB) = 2^{12}
- Number of page table entries = $2^{32} / 2^{12} = 2^{20}$
- Number of pages in a 4GB DRAM = $2^{32} / 2^{12} = 2^{20}$
- Size of each page table entry =
20 bits + page table bits ~ 3 bytes

Virtual page number	Physical page number
7	3
6	x
5	x
4	x
3	x
2	4
1	6
0	0

Page table size

- Virtual addresses: 2^{32}
- Size of page = (4KB) = 2^{12}
- Number of page table entries = $2^{32} / 2^{12} = 2^{20}$
- Number of pages in a 4GB DRAM = $2^{32} / 2^{12} = 2^{20}$
- Size of each page table entry =
20 bits + page table bits ~ 3 bytes
- Size of page table = $3 * 2^{20} = 3\text{MB!}$

Virtual page number	Physical page number
7	3
6	x
5	x
4	x
3	x
2	4
1	6
0	0

Page table size

- Virtual addresses: 2^{32}
- Size of page = (4KB) = 2^{12}
- Number of page table entries = $2^{32} / 2^{12} = 2^{20}$
- Number of pages in a 4GB DRAM = $2^{32} / 2^{12} = 2^{20}$
- Size of each page table entry =
20 bits + page table bits ~ 3 bytes
- Size of page table = $3 * 2^{20} = 3\text{MB}$!
- 1024 processes => ~3GB memory!

Virtual page number	Physical page number
7	3
6	x
5	x
4	x
3	x
2	4
1	6
0	0

Reducing page table size

- Bigger pages
 - Virtual addresses: 2^{32}
 - Size of page = (**4MB**) = 2^{22}

Reducing page table size

- Bigger pages
 - Virtual addresses: 2^{32}
 - Size of page = (**4MB**) = 2^{22}
 - Number of page table entries = $2^{32} / 2^{22} = 2^{10}$

Reducing page table size

- Bigger pages
 - Virtual addresses: 2^{32}
 - Size of page = (**4MB**) = 2^{22}
 - Number of page table entries = $2^{32} / 2^{22} = 2^{10}$
 - Number of pages in a **4GB** DRAM = $2^{32} / 2^{22} = 2^{10}$

Reducing page table size


- Bigger pages
 - Virtual addresses: 2^{32}
 - Size of page = **(4MB)** = 2^{22}
 - Number of page table entries = $2^{32} / 2^{22} = 2^{10}$
 - Number of pages in a **4GB** DRAM = $2^{32} / 2^{22} = 2^{10}$
 - Size of each page table entry = **10** bits ~ **2 bytes**

Reducing page table size

- Bigger pages
 - Virtual addresses: 2^{32}
 - Size of page = **(4MB)** = 2^{22}
 - Number of page table entries = $2^{32} / 2^{22} = 2^{10}$
 - Number of pages in a **4GB** DRAM = $2^{32} / 2^{22} = 2^{10}$
 - Size of each page table entry = **10** bits ~ **2 bytes**
 - Size of page table = $2 * 2^{10} = 2\text{KB!}$


Reducing page table size

- Bigger pages
 - Virtual addresses: 2^{32}
 - Size of page = **(4MB)** = 2^{22}
 - Number of page table entries = $2^{32} / 2^{22} = 2^{10}$
 - Number of pages in a **4GB** DRAM = $2^{32} / 2^{22} = 2^{10}$
 - Size of each page table entry = **10** bits ~ **2 bytes**
 - Size of page table = $2 * 2^{10} = 2\text{KB}$!
- Bigger pages increase internal fragmentation



Observation


- Lot of page table entries are invalid



Virtual page number	Physical page number
7	3
6	x
5	x
4	x
3	x
2	4
1	6
0	0

Observation

- Lot of page table entries are invalid



Process 1 address space

Virtual page number	Physical page number
7	3
6	x
5	x
4	x
3	x
2	4
1	6
0	0

Multi-level page table

Virtual page number	Physical page number
9	3
8	1
7	x
6	x
5	x
4	x
3	x
2	x
1	6
0	0

PPN: 8

Virtual page number	Physical page number
8,9	9
6,7	x
4,5	x
2,3	x
0,1	10

PPN: 9

Virtual page number	Physical page number
9	3
8	1

PPN: 10

Virtual page number	Physical page number
1	6
0	0

Multi-level page table

Virtual page number	Physical page number
9	3
8	1
7	x
6	x
5	x
4	x
3	x
2	x
1	6
0	0

PPN: 8

Virtual page number	Physical page number
8,9	9
6,7	x
4,5	x
2,3	x
0,1	10

PPN: 9

Virtual page number	Physical page number
9	3
8	1

PPN: 10

Virtual page number	Physical page number
1	6
0	0

- Page directory entries point to page table pages
- Unused portions of virtual address space are skipped!

Notebook analogy


Paging



0	1	2	3	4	5
xv6	is	an	OS	for	x86




0	1	2	3
Write	an	SQL	query



0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
OS:1 DB:5	0:3,1:4, 2:2, 3:8 , 4:6,5:10	an	xv6	is	0:9,1:2, 2:7,3:12	for	SQL	OS	Write	x86		query		

Reading **S**

- Read second letter from 3rd page




- Read second letter from 8th page

Notebook analogy

Page directories: call 4 pages a “section”

Reading S

- Read second letter from 3rd page in **Section 0**



- Read second letter from 8th page



OS

0	1	2	3	0	1	2	3
xv6	is	an	OS	for	x86		
Section 0				Section 1			



DB

0	1	2	3
Write	an	SQL	query
Section 0			



Notebook

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14
OS:1 DB:5	0:11, 1:14	an	xv6	is	0:13	for	SQL	OS	Write	x86	0:3,1:4, 2:2, 3:8	query	0:9,1:2, 2:7,3:12	0:6, 1:10


Address translation

Simple address space

- 16KB address space has 2^{14} addresses
- Each page has 64 ($= 2^6$) bytes
- Number of pages $= 2^{14} / 2^6 = 2^8$
- First 8 bits are page number, last 6 bits are offset within the page


0000 0000	code
0000 0001	code
0000 0010	(free)
0000 0011	(free)
0000 0100	heap
0000 0101	heap
0000 0110	(free)
0000 0111	(free)
.....	... all free ...
1111 1100	(free)
1111 1101	(free)
1111 1110	stack
1111 1111	stack

Figure 20.4: A 16KB Address Space With 64-byte Pages




Address translation

Simple address space



Address translation

Simple address space




Page Directory		Page of PT (@PFN:100)			Page of PT (@PFN:101)		
PFN	valid?	PFN	valid	prot	PFN	valid	prot
100	1	10	1	r-x	—	0	—
—	0	23	1	r-x	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	80	1	rw-	—	0	—
—	0	59	1	rw-	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	55	1	rw-
101	1	—	0	—	45	1	rw-

Figure 20.5: A Page Directory, And Pieces Of Page Table

Address translation

Simple address space



- Example: 0x3F81 (VA) => 0x0DC1 (PA)

Page Directory		Page of PT (@PFN:100)			Page of PT (@PFN:101)		
PFN	valid?	PFN	valid	prot	PFN	valid	prot
100	1	10	1	r-x	—	0	—
—	0	23	1	r-x	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	80	1	rw-	—	0	—
—	0	59	1	rw-	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	55	1	rw-
101	1	—	0	—	45	1	rw-


Figure 20.5: A Page Directory, And Pieces Of Page Table

Address translation

Simple address space

- Example: 0x3F81 (VA) => 0x0DC1 (PA)

Page dir idx 15				Page table idx 14				Offset					
1	1	1	1	1	1	1	0	0	0	0	0	0	1
3		F				8				1			



Page Directory		Page of PT (@PFN:100)			Page of PT (@PFN:101)		
PFN	valid?	PFN	valid	prot	PFN	valid	prot
100	1	10	1	r-x	—	0	—
—	0	23	1	r-x	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	80	1	rw-	—	0	—
—	0	59	1	rw-	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	55	1	rw-
101	1	—	0	—	45	1	rw-


Figure 20.5: A Page Directory, And Pieces Of Page Table

Address translation

Simple address space

- Example: 0x3F81 (VA) => 0x0DC1 (PA)

Page dir idx 15				Page table idx 14				Offset					
1	1	1	1	1	1	1	0	0	0	0	0	0	1
3		F				8				1			



Page Directory		Page of PT (@PFN:100)			Page of PT (@PFN:101)		
PFN	valid?	PFN	valid	prot	PFN	valid	prot
100	1	10	1	r-x	—	0	—
—	0	23	1	r-x	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	80	1	rw-	—	0	—
—	0	59	1	rw-	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	55	1	rw-
15 101	1	—	0	—	45	1	rw-


Figure 20.5: A Page Directory, And Pieces Of Page Table

Address translation

Simple address space

- Example: 0x3F81 (VA) => 0x0DC1 (PA)

Page dir idx 15				Page table idx 14				Offset					
1	1	1	1	1	1	1	0	0	0	0	0	0	1
3		F				8				1			




Page Directory		Page of PT (@PFN:100)			Page of PT (@PFN:101)		
PFN	valid?	PFN	valid	prot	PFN	valid	prot
100	1	10	1	r-x	—	0	—
—	0	23	1	r-x	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	80	1	rw-	—	0	—
—	0	59	1	rw-	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	55	1	rw-
15	1	101	1	—	45	1	rw-

Figure 20.5: A Page Directory, And Pieces Of Page Table

Address translation

Simple address space




- Example: 0x3F81 (VA) => 0x0DC1 (PA)

Page dir idx 15				Page table idx 14				Offset					
1	1	1	1	1	1	1	0	0	0	0	0	0	1
3		F				8		1					

Page Directory		Page of PT (@PFN:100)			Page of PT (@PFN:101)		
PFN	valid?	PFN	valid	prot	PFN	valid	prot
100	1	10	1	r-x	—	0	—
—	0	23	1	r-x	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	80	1	rw-	—	0	—
—	0	59	1	rw-	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	55	1	rw-
15	1	101	1	—	45	1	rw-

Figure 20.5: A Page Directory, And Pieces Of Page Table

Simple address space



Physical page number 55								Offset					
0	0	1	1	0	1	1	1	0	0	0	0	0	1
0		D			C			1					

Figure 20.5: A Page Directory, And Pieces Of Page Table

x86 segmentation and paging

- Segmentation:
 - Virtual address (logical address) => “linear address”
- Paging:
 - Linear address => Physical address




Figure 3-1. Segmentation and Paging

Address translation with paging on x86





Figure 3-12. Linear Address Translation (4-KByte Pages)

Address translation with paging on x86



x86 PTEs, PDEs

- 2^{20} 4KB pages in a 4GB DRAM
- Page base address, page table base address are 20 bits





Figure 3-14. Format of Page-Directory and Page-Table Entries for 4-KByte Pages and 32-Bit Physical Addresses

x86 PTEs, PDEs

- 2^{20} 4KB pages in a 4GB DRAM
 - Page base address, page table base address are 20 bits
- Bits set by OS, used by hardware:





Figure 3-14. Format of Page-Directory and Page-Table Entries for 4-KByte Pages and 32-Bit Physical Addresses

x86 PTEs, PDEs

- 2^{20} 4KB pages in a 4GB DRAM
 - Page base address, page table base address are 20 bits
- Bits set by OS, used by hardware:
 - **Present:** It is a valid entry




Figure 3-14. Format of Page-Directory and Page-Table Entries for 4-KByte Pages and 32-Bit Physical Addresses

x86 PTEs, PDEs

- 2^{20} 4KB pages in a 4GB DRAM
 - Page base address, page table base address are 20 bits
- Bits set by OS, used by hardware:
 - **Present**: It is a valid entry
 - **Read/write**: Can write if 1


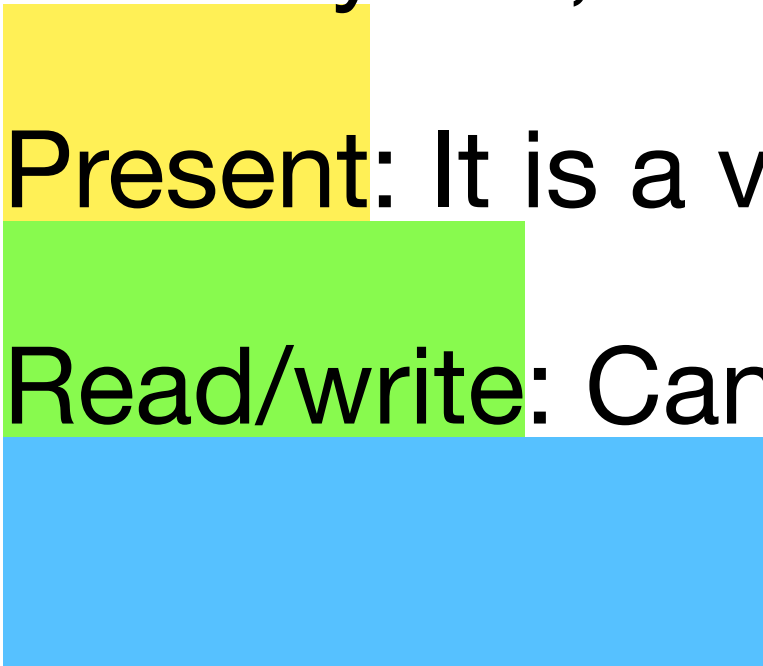


Figure 3-14. Format of Page-Directory and Page-Table Entries for 4-KByte Pages and 32-Bit Physical Addresses

x86 PTEs, PDEs

- 2^{20} 4KB pages in a 4GB DRAM
 - Page base address, page table base address are 20 bits
- Bits set by OS, used by hardware:
 - **Present**: It is a valid entry
 - **Read/write**: Can write if 1
 - **User/supervisor**: CPL=3 can access if 1




Figure 3-14. Format of Page-Directory and Page-Table Entries for 4-KByte Pages and 32-Bit Physical Addresses

x86 PTEs, PDEs

- 2^{20} 4KB pages in a 4GB DRAM
 - Page base address, page table base address are 20 bits
- Bits set by OS, used by hardware:
 - **Present**: It is a valid entry
 - **Read/write**: Can write if 1
 - **User/supervisor**: CPL=3 can access if 1
- Bits set by hardware, used/cleared by OS:
 - **Available for system programmer's use**
 - **Global Page**
 - **Page Table Attribute Index**
 - **Dirty**
 - **Accessed**
 - **Cache Disabled**
 - **Write-Through**
 - **User/Supervisor**
 - **Read/Write**
 - **Present**




Figure 3-14. Format of Page-Directory and Page-Table Entries for 4-KByte Pages and 32-Bit Physical Addresses

x86 PTEs, PDEs

- 2^{20} 4KB pages in a 4GB DRAM
 - Page base address, page table base address are 20 bits
- Bits set by OS, used by hardware:
 - **Present**: It is a valid entry
 - **Read/write**: Can write if 1
 - **User/supervisor**: CPL=3 can access if 1
- Bits set by hardware, used/cleared by OS:
 - **Accessed**: Hardware accessed this page




Figure 3-14. Format of Page-Directory and Page-Table Entries for 4-KByte Pages and 32-Bit Physical Addresses

x86 PTEs, PDEs

- 2^{20} 4KB pages in a 4GB DRAM
 - Page base address, page table base address are 20 bits
- Bits set by OS, used by hardware:
 - **Present**: It is a valid entry
 - **Read/write**: Can write if 1
 - **User/supervisor**: CPL=3 can access if 1
- Bits set by hardware, used/cleared by OS:
 - **Accessed**: Hardware accessed this page
 - **Dirty**: Hardware wrote to this page







Figure 3-14. Format of Page-Directory and Page-Table Entries for 4-KByte Pages and 32-Bit Physical Addresses

Performance degradation!



- Accessing 1 memory location requires accessing 3 memory locations!



	Page Directory		Page of PT (@PFN:100)			Page of PT (@PFN:101)		
	PFN	valid?	PFN	valid	prot	PFN	valid	prot
	100	1	10	1	r-x	—	0	—
	—	0	23	1	r-x	—	0	—
	—	0	—	0	—	—	0	—
	—	0	—	0	—	—	0	—
	—	0	80	1	rw-	—	0	—
	—	0	59	1	rw-	—	0	—
	—	0	—	0	—	—	0	—
	—	0	—	0	—	—	0	—
	—	0	—	0	—	—	0	—
	—	0	—	0	—	—	0	—
	—	0	—	0	—	—	0	—
	—	0	—	0	—	—	0	—
	—	0	—	0	—	—	0	—
	—	0	—	0	—	—	0	—
	—	0	—	0	—	—	0	—
	—	0	—	0	—	—	0	—
	—	0	—	0	—	55	1	rw-
→ 15	101	1	—	0	—	45	1	rw-

Figure 20.5: A Page Directory, And Pieces Of Page Table

Translation-lookaside buffer (TLB)

- First check page translation in TLB before walking the page table
- TLB hit: ~0.5-1 cycle
- TLB miss: ~10-100 cycles

Page dir idx 15				Page table idx 14				Offset					
1	1	1	1	1	1	1	0	0	0	0	0	0	1
3		F				8			1				

Physical page number 55								Offset					
0	0	1	1	0	1	1	1	0	0	0	0	0	1
0		D				C			1				

TLB

VPN								PPN						
1	1	1	1	1	1	1	0	0	0	1	1	0	1	1

Which programs will run faster?

Which programs will have lesser TLB misses?

```
int sum = 0;
for (i = 0; i < 10; i++) {
    sum += a[i];
}
```

	Offset				
	00	04	08	12	16
VPN = 00					
VPN = 01					
VPN = 02					
VPN = 03					
VPN = 04					
VPN = 05					
VPN = 06		a[0]	a[1]	a[2]	
VPN = 07	a[3]	a[4]	a[5]	a[6]	
VPN = 08	a[7]	a[8]	a[9]		
VPN = 09					
VPN = 10					
VPN = 11					
VPN = 12					
VPN = 13					
VPN = 14					
VPN = 15					

Which programs will run faster?

Which programs will have lesser TLB misses?

```
int sum = 0;
for (i = 0; i < 10; i++) {
    sum += a[i];
}
```

	Offset				
	00	04	08	12	16
VPN = 00					
VPN = 01					
VPN = 02					
VPN = 03					
VPN = 04					
VPN = 05					
VPN = 06		a[0]	a[1]	a[2]	
VPN = 07	a[3]	a[4]	a[5]	a[6]	
VPN = 08	a[7]	a[8]	a[9]		
VPN = 09					
VPN = 10					
VPN = 11					
VPN = 12					
VPN = 13					
VPN = 14					
VPN = 15					

Which programs will run faster?

Which programs will have lesser TLB misses?

```
int sum = 0;
for (i = 0; i < 10; i++) {
    sum += a[i];
}
```

	Offset				
	00	04	08	12	16
VPN = 00					
VPN = 01					
VPN = 02					
VPN = 03					
VPN = 04					
VPN = 05					
VPN = 06		a[0]	a[1]	a[2]	
VPN = 07	a[3]	a[4]	a[5]	a[6]	
VPN = 08	a[7]	a[8]	a[9]		
VPN = 09					
VPN = 10					
VPN = 11					
VPN = 12					
VPN = 13					
VPN = 14					
VPN = 15					

Which programs will run faster?

Which programs will have lesser TLB misses?

```
int sum = 0;
for (i = 0; i < 10; i++) {
    sum += a[i];
}
```

	Offset				
	00	04	08	12	16
VPN = 00					
VPN = 01					
VPN = 02					
VPN = 03					
VPN = 04					
VPN = 05					
VPN = 06		a[0]	a[1]	a[2]	
VPN = 07	a[3]	a[4]	a[5]	a[6]	
VPN = 08	a[7]	a[8]	a[9]		
VPN = 09					
VPN = 10					
VPN = 11					
VPN = 12					
VPN = 13					
VPN = 14					
VPN = 15					

Which programs will run faster?

Which programs will have lesser TLB misses?

```
int sum = 0;
for (i = 0; i < 10; i++) {
    sum += a[i];
}
```

	Offset				
	00	04	08	12	16
VPN = 00					
VPN = 01					
VPN = 02					
VPN = 03					
VPN = 04					
VPN = 05					
VPN = 06	✗		a[0]	✓ a[1]	✓ a[2]
VPN = 07		a[3]	a[4]	a[5]	a[6]
VPN = 08		a[7]	a[8]	a[9]	
VPN = 09					
VPN = 10					
VPN = 11					
VPN = 12					
VPN = 13					
VPN = 14					
VPN = 15					

Which programs will run faster?

Which programs will have lesser TLB misses?

```
int sum = 0;
for (i = 0; i < 10; i++) {
    sum += a[i];
}
```

	Offset				
	00	04	08	12	16
VPN = 00					
VPN = 01					
VPN = 02					
VPN = 03					
VPN = 04					
VPN = 05					
VPN = 06		a[0]	a[1]	a[2]	
VPN = 07	a[3]	a[4]	a[5]	a[6]	
VPN = 08	a[7]	a[8]	a[9]		
VPN = 09					
VPN = 10					
VPN = 11					
VPN = 12					
VPN = 13					
VPN = 14					
VPN = 15					

Which programs will run faster?

Which programs will have lesser TLB misses?

```
int sum = 0;
for (i = 0; i < 10; i++) {
    sum += a[i];
}
```

Offset				
00	04	08	12	16
VPN = 00				
VPN = 01				
VPN = 02				
VPN = 03				
VPN = 04				
VPN = 05				
VPN = 06		a[0]	a[1]	a[2]
VPN = 07	a[3]	a[4]	a[5]	a[6]
VPN = 08	a[7]	a[8]	a[9]	
VPN = 09				
VPN = 10				
VPN = 11				
VPN = 12				
VPN = 13				
VPN = 14				
VPN = 15				

Which programs will run faster?

Which programs will have lesser TLB misses?

```
int sum = 0;
for (i = 0; i < 10; i++) {
    sum += a[i];
}
```

	Offset				
	00	04	08	12	16
VPN = 00					
VPN = 01					
VPN = 02					
VPN = 03					
VPN = 04					
VPN = 05					
VPN = 06		a[0]	a[1]	a[2]	
VPN = 07	a[3]	a[4]	a[5]	a[6]	
VPN = 08	a[7]	a[8]	a[9]		
VPN = 09					
VPN = 10					
VPN = 11					
VPN = 12					
VPN = 13					
VPN = 14					
VPN = 15					

Which programs will run faster?

Which programs will have lesser TLB misses?

```
int sum = 0;
for (i = 0; i < 10; i++) {
    sum += a[i];
}
```

Offset				
00	04	08	12	16
VPN = 00				
VPN = 01				
VPN = 02				
VPN = 03				
VPN = 04				
VPN = 05				
VPN = 06		a[0]	a[1]	a[2]
VPN = 07	a[3]	a[4]	a[5]	a[6]
VPN = 08	a[7]	a[8]	a[9]	
VPN = 09				
VPN = 10				
VPN = 11				
VPN = 12				
VPN = 13				
VPN = 14				
VPN = 15				

Which programs will run faster?

Which programs will have lesser TLB misses?

```
int sum = 0;
for (i = 0; i < 10; i++) {
    sum += a[i];
}
```

		Offset				
		00	04	08	12	16
VPN = 00						
VPN = 01						
VPN = 02						
VPN = 03						
VPN = 04						
VPN = 05						
VPN = 06			a[0]	a[1]	a[2]	
VPN = 07	a[3]	a[4]	a[5]	a[6]		
VPN = 08	a[7]	a[8]	a[9]			
VPN = 09						
VPN = 10						
VPN = 11						
VPN = 12						
VPN = 13						
VPN = 14						
VPN = 15						

Which programs will run faster?

Which programs will have lesser TLB misses?

```
int sum = 0;
for (i = 0; i < 10; i++) {
    sum += a[i];
}
```

	Offset				
	00	04	08	12	16
VPN = 00					
VPN = 01					
VPN = 02					
VPN = 03					
VPN = 04					
VPN = 05					
VPN = 06		a[0]	a[1]	a[2]	
VPN = 07	a[3]	a[4]	a[5]	a[6]	
VPN = 08	a[7]	a[8]	a[9]		
VPN = 09					
VPN = 10					
VPN = 11					
VPN = 12					
VPN = 13					
VPN = 14					
VPN = 15					

Which programs will run faster?

Which programs will have lesser TLB misses?

```
int sum = 0;
for (i = 0; i < 10; i++) {
    sum += a[i];
}
```

	Offset				
	00	04	08	12	16
VPN = 00					
VPN = 01					
VPN = 02					
VPN = 03					
VPN = 04					
VPN = 05					
VPN = 06		a[0]	a[1]	a[2]	
VPN = 07	a[3]	a[4]	a[5]	a[6]	
VPN = 08	a[7]	a[8]	a[9]		
VPN = 09					
VPN = 10					
VPN = 11					
VPN = 12					
VPN = 13					
VPN = 14					
VPN = 15					

- High **spatial locality**: after the program accessed a memory location, it will access a nearby memory location

Which programs will run faster?

Which programs will have lesser TLB misses?

```
int sum = 0;
for (i = 0; i < 10; i++) {
    sum += a[i];
}
```

	Offset				
	00	04	08	12	16
VPN = 00					
VPN = 01					
VPN = 02					
VPN = 03					
VPN = 04					
VPN = 05					
VPN = 06		a[0]	a[1]	a[2]	
VPN = 07	a[3]	a[4]	a[5]	a[6]	
VPN = 08	a[7]	a[8]	a[9]		
VPN = 09					
VPN = 10					
VPN = 11					
VPN = 12					
VPN = 13					
VPN = 14					
VPN = 15					

- High **spatial locality**: after the program accessed a memory location, it will access a nearby memory location
- High **temporal locality**: after the program accessed a memory location, it will soon access it again

Example bad programs

- Low spatial and temporal locality: most accesses lead to TLB miss

Example bad programs

- Low spatial and temporal locality: most accesses lead to TLB miss
- Large hash table with random access

Example bad programs

- Low spatial and temporal locality: most accesses lead to TLB miss
- Large hash table with random access

```
int get(int I)
    return a[I];
```

Example bad programs

- Low spatial and temporal locality: most accesses lead to TLB miss
- Large hash table with random access

```
int get(int I)
    return a[I];
```

VPN=1	a[0]	a[1]	a[2]	a[3]
VPN=2	a[4]	a[5]	a[6]	a[7]
VPN=3	a[8]	a[9]	a[10]	a[11]
VPN=4	a[12]	a[13]	a[14]	a[15]
VPN=5	a[16]	a[17]	a[18]	a[19]

Example bad programs

- Low spatial and temporal locality: most accesses lead to TLB miss

- Large hash table with random access

```
int get(int I)
```

```
    return a[I];
```

VPN=1	a[0]	a[1]	a[2]	a[3]
VPN=2	a[4]	a[5]	a[6]	a[7]
VPN=3	a[8]	a[9]	a[10]	a[11]
VPN=4	a[12]	a[13]	a[14]	a[15]
VPN=5	a[16]	a[17]	a[18]	a[19]

- Traversing large graphs: two neighbours can be on different pages

Example bad programs

- Low spatial and temporal locality: most accesses lead to TLB miss

- Large hash table with random access

```
int get(int I)
```

```
    return a[I];
```

VPN=1	a[0]	a[1]	a[2]	a[3]
VPN=2	a[4]	a[5]	a[6]	a[7]
VPN=3	a[8]	a[9]	a[10]	a[11]
VPN=4	a[12]	a[13]	a[14]	a[15]
VPN=5	a[16]	a[17]	a[18]	a[19]

- Traversing large graphs: two neighbours can be on different pages
- Ok for small hash tables, small graphs.
 - Working set of the program: Amount of memory that it touches

Example bad programs

- Low spatial and temporal locality: most accesses lead to TLB miss

- Large hash table with random access

```
int get(int I)
```

```
    return a[I];
```

VPN=1	a[0]	a[1]	a[2]	a[3]
VPN=2	a[4]	a[5]	a[6]	a[7]
VPN=3	a[8]	a[9]	a[10]	a[11]
VPN=4	a[12]	a[13]	a[14]	a[15]
VPN=5	a[16]	a[17]	a[18]	a[19]

- Traversing large graphs: two neighbours can be on different pages
- Ok for small hash tables, small graphs.
 - Working set of the program: Amount of memory that it touches
- $\text{TLB reach} = (\text{number of TLB entries}) * (\text{page size})$

Example bad programs

- Low spatial and temporal locality: most accesses lead to TLB miss

- Large hash table with random access

```
int get(int I)
```

```
    return a[I];
```

VPN=1	a[0]	a[1]	a[2]	a[3]
VPN=2	a[4]	a[5]	a[6]	a[7]
VPN=3	a[8]	a[9]	a[10]	a[11]
VPN=4	a[12]	a[13]	a[14]	a[15]
VPN=5	a[16]	a[17]	a[18]	a[19]


- Traversing large graphs: two neighbours can be on different pages
- Ok for small hash tables, small graphs.
 - Working set of the program: Amount of memory that it touches
- TLB reach = (number of TLB entries) * (page size)
- Bad performance when Working set > TLB reach

Increasing TLB reach

- Larger TLBs
 - 32 -> 64 entries.
 - Larger caches => slower hits

Increasing TLB reach


- Larger TLBs
 - 32 -> 64 entries.
 - Larger caches => slower hits
- Larger pages
 - 1 TLB entry for 4MB of addresses (with a 4MB page)
 - 1024 TLB entries for 4MB of addresses (with 4KB pages)



Increasing TLB reach

- Larger TLBs
 - 32 -> 64 entries.
 - Larger caches => slower hits
- Larger pages
 - 1 TLB entry for 4MB of addresses (with a 4MB page)
 - 1024 TLB entries for 4MB of addresses (with 4KB pages)
- Allocating large pages on Linux:

```
posix_memalign(void **memptr, size_t alignment, size_t size);  
int madvise(void *addr, size_t length, int advice= MADV_HUGEPAGE);
```



TLB on my x86-64 machine

```
$ cpuid
```

```
..  
(simple synth) = Intel Core (unknown type) (Kaby Lake / Coffee Lake) {Skylake}, 14nm
```

```
..  
cache and TLB information (2):  
  0x63: data TLB: 2M/4M pages, 4-way, 32 entries  
        data TLB: 1G pages, 4-way, 4 entries  
  0x03: data TLB: 4K pages, 4-way, 64 entries  
  0x76: instruction TLB: 2M/4M pages, fully, 8 entries  
  0xb5: instruction TLB: 4K, 8-way, 64 entries  
  0xc3: L2 TLB: 4K/2M pages, 6-way, 1536 entries
```

```
..
```

TLB on my x86-64 machine

```
$ cpuid
```

```
..  
(simple synth) = Intel Core (unknown type) (Kaby Lake / Coffee Lake) {Skylake}, 14nm
```

```
..  
cache and TLB information (2):
```

```
0x63: data TLB: 2M/4M pages, 4-way, 32 entries
```

```
data TLB: 1G pages, 4-way, 4 entries
```

```
0x03: data TLB: 4K pages, 4-way, 64 entries
```


```
0x76: instruction TLB: 2M/4M pages, fully, 8 entries
```

```
0xb5: instruction TLB: 4K, 8-way, 64 entries
```

```
0xc3: L2 TLB: 4K/2M pages, 6-way, 1536 entries
```

```
..
```

Separate TLBs for different page sizes



TLB on my x86-64 machine

```
$ cpuid
```

```
..  
(simple synth) = Intel Core (unknown type) (Kaby Lake / Coffee Lake) {Skylake}, 14nm
```

```
..  
cache and TLB information (2):
```

```
0x63: data TLB: 2M/4M pages, 4-way, 32 entries
```

```
data TLB: 1G pages, 4-way, 4 entries
```

```
0x03: data TLB: 4K pages, 4-way, 64 entries
```


```
0x76: instruction TLB: 2M/4M pages, fully, 8 entries
```

```
0xb5: instruction TLB: 4K, 8-way, 64 entries
```

```
0xc3: L2 TLB: 4K/2M pages, 6-way, 1536 entries
```

```
..
```

Separate TLBs for different page sizes



Data TLB reach:

- $4\text{KB} * 64 = 256\text{KB}$
- $4\text{MB} * 32 = 128\text{MB}$
- $1\text{GB} * 4 = 4\text{GB}$

TLB on my x86-64 machine

```
$ cpuid
```

```
..  
(simple synth) = Intel Core (unknown type) (Kaby Lake / Coffee Lake) {Skylake}, 14nm
```

```
..  
cache and TLB information (2):
```

```
0x63: data TLB: 2M/4M pages, 4-way, 32 entries
```

```
data TLB: 1G pages, 4-way, 4 entries
```

```
0x03: data TLB: 4K pages, 4-way, 64 entries
```

```
0x76: instruction TLB: 2M/4M pages, fully, 8 entries
```

```
0xb5: instruction TLB: 4K, 8-way, 64 entries
```

```
0xc3: L2 TLB: 4K/2M pages, 6-way, 1536 entries
```

```
..
```

Separate TLBs for different page sizes

Separate TLBs for instruction and data

Data TLB reach:

- $4\text{KB} * 64 = 256\text{KB}$
- $4\text{MB} * 32 = 128\text{MB}$
- $1\text{GB} * 4 = 4\text{GB}$

TLB on my x86-64 machine

```
$ cpuid
```

```
..  
(simple synth) = Intel Core (unknown type) (Kaby Lake / Coffee Lake) {Skylake}, 14nm
```

```
..  
cache and TLB information (2):
```

```
0x63: data TLB: 2M/4M pages, 4-way, 32 entries
```

```
data TLB: 1G pages, 4-way, 4 entries
```

```
0x03: data TLB: 4K pages, 4-way, 64 entries
```

```
0x76: instruction TLB: 2M/4M pages, fully, 8 entries
```

```
0xb5: instruction TLB: 4K, 8-way, 64 entries
```

```
0xc3: L2 TLB: 4K/2M pages, 6-way, 1536 entries
```

```
..
```

Separate TLBs for different page sizes

Separate TLBs for instruction and data

Data TLB reach:

- $4\text{KB} * 64 = 256\text{KB}$
- $4\text{MB} * 32 = 128\text{MB}$
- $1\text{GB} * 4 = 4\text{GB}$

Instruction TLB reach:

- $4\text{KB} * 64 = 256\text{KB}$
- $4\text{MB} * 8 = 32\text{MB}$

TLB on my x86-64 machine

```
$ cpuid
```

```
..  
(simple synth) = Intel Core (unknown type) (Kaby Lake / Coffee Lake) {Skylake}, 14nm
```

```
..  
cache and TLB information (2):
```

```
0x63: data TLB: 2M/4M pages, 4-way, 32 entries
```

```
data TLB: 1G pages, 4-way, 4 entries
```

```
0x03: data TLB: 4K pages, 4-way, 64 entries
```

```
0x76: instruction TLB: 2M/4M pages, fully, 8 entries
```

```
0xb5: instruction TLB: 4K, 8-way, 64 entries
```

```
0xc3: L2 TLB: 4K/2M pages, 6-way, 1536 entries
```

```
..
```

Separate TLBs for different page sizes

Separate TLBs for instruction and data

Much larger (slower) L2 TLB

Data TLB reach:

- $4\text{KB} * 64 = 256\text{KB}$
- $4\text{MB} * 32 = 128\text{MB}$
- $1\text{GB} * 4 = 4\text{GB}$

Instruction TLB reach:

- $4\text{KB} * 64 = 256\text{KB}$
- $4\text{MB} * 8 = 32\text{MB}$

TLB on my x86-64 machine

```
$ cpuid
```

```
..  
(simple synth) = Intel Core (unknown type) (Kaby Lake / Coffee Lake) {Skylake}, 14nm
```

```
..  
cache and TLB information (2):
```

```
0x63: data TLB: 2M/4M pages, 4-way, 32 entries
```

```
data TLB: 1G pages, 4-way, 4 entries
```

```
0x03: data TLB: 4K pages, 4-way, 64 entries
```

```
0x76: instruction TLB: 2M/4M pages, fully, 8 entries
```

```
0xb5: instruction TLB: 4K, 8-way, 64 entries
```

```
0xc3: L2 TLB: 4K/2M pages, 6-way, 1536 entries
```

```
..
```

Separate TLBs for different page sizes

Separate TLBs for instruction and data

Much larger (slower) L2 TLB

Data TLB reach:

- $4\text{KB} * 64 = 256\text{KB}$
- $4\text{MB} * 32 = 128\text{MB}$
- $1\text{GB} * 4 = 4\text{GB}$

Instruction TLB reach:

- $4\text{KB} * 64 = 256\text{KB}$
- $4\text{MB} * 8 = 32\text{MB}$

L2 TLB reach:


- $4\text{KB} * 1536 = 6\text{MB}$
- $2\text{MB} * 1536 = 3\text{GB}$

TLB size, multiple TLBs example

```
int jump = PAGE_SIZE / sizeof(int);  
for (i = 0; i < NUM_PAGES * jump; i += jump) {  
    a[i] += 1;  
}
```


TLB size, multiple TLBs example

```
int jump = PAGE_SIZE / sizeof(int);  
for (i = 0; i < NUMPAGES * jump; i += jump) {  
    a[i] += 1;  
}
```




TLB size, multiple TLBs example

```
int jump = PAGE_SIZE / sizeof(int);  
for (i = 0; i < NUMPAGES * jump; i += jump) {  
    a[i] += 1;  
}
```




TLB size, multiple TLBs example

```
int jump = PAGE_SIZE / sizeof(int);  
for (i = 0; i < NUMPAGES * jump; i += jump) {  
    a[i] += 1;  
}
```



TLB size, multiple TLBs example

```
int jump = PAGE_SIZE / sizeof(int);  
for (i = 0; i < NUMPAGES * jump; i += jump) {  
    a[i] += 1;  
}
```



TLB replacement policies

- Need to replace an entry after TLB is full. Which entry to replace to minimise TLB miss rate?

TLB replacement policies

- Need to replace an entry after TLB is full. Which entry to replace to minimise TLB miss rate?
- Least recently used (LRU):
 - If an entry hasn't been used recently, it is unlikely to be used soon => assumes spatial and temporal locality of access.
 - Corner case behaviours: Program cycling over $N+1$ pages.

TLB replacement policies

- Need to replace an entry after TLB is full. Which entry to replace to minimise TLB miss rate?
- Least recently used (LRU):
 - If an entry hasn't been used recently, it is unlikely to be used soon => assumes spatial and temporal locality of access.
 - Corner case behaviours: Program cycling over $N+1$ pages.
- Hardware typically implements something simple
 - FIFO
 - Random replacement: Just pick an entry randomly


Context switching

- During context switch, OS changes CR3 register to change page table
- Privileged operation

Virtual page number	Physical page number
7	3
6	x
5	x
4	x
3	x
2	4
1	6
0	0

Virtual page number	Physical page number
7	5
6	x
5	x
4	x
3	7
2	3
1	6
0	1

Virtual page number	Physical page number
0	0
2	4




Context switching

- During context switch, OS changes CR3 register to change page table
- Privileged operation
- Marks each TLB entry invalid. Every memory access after context switch causes TLB miss!

Virtual page number	Physical page number
0	0
2	4

Virtual page number	Physical page number
7	3
6	x
5	x
4	x
3	x
2	4
1	6
0	0

Virtual page number	Physical page number
7	5
6	x
5	x
4	x
3	7
2	3
1	6
0	1




Context switching

- During context switch, OS changes CR3 register to change page table
- Privileged operation
- Marks each TLB entry invalid. Every memory access after context switch causes TLB miss!

Virtual page number	Physical page number
0	0
2	4

Virtual page number	Physical page number
7	3
6	x
5	x
4	x
3	x
2	4
1	6
0	0

Virtual page number	Physical page number
7	5
6	x
5	x
4	x
3	7
2	3
1	6
0	1




Context switching

- During context switch, OS changes CR3 register to change page table
- Privileged operation
- Marks each TLB entry invalid. Every memory access after context switch causes TLB miss!

Virtual page number	Physical page number
0	0
2	4

Virtual page number	Physical page number
7	3
6	x
5	x
4	x
3	x
2	4
1	6
0	0

Virtual page number	Physical page number
7	5
6	x
5	x
4	x
3	7
2	3
1	6
0	1




Context switching

- During context switch, OS changes CR3 register to change page table
- Privileged operation
- Marks each TLB entry invalid. Every memory access after context switch causes TLB miss!

Virtual page number	Physical page number
7	3
6	x
5	x
4	x
3	x
2	4
1	6
0	0

Virtual page number	Physical page number
7	5
6	x
5	x
4	x
3	7
2	3
1	6
0	1

Virtual page number	Physical page number
0	0
2	4
0	1



INVLPG instruction

- TLB is neither write-back, nor write-through cache
 - Need to run INVLPG <virtual page number> when a page table entry is modified
- Similar to how OS needs to reload segment selector when corresponding GDT entry is changed

Tagged TLBs

- TLB entries are “tagged” with a *process context identifier*

TLB

Virtual page number	Physical page number	Process context identifier
0	0	P1
2	4	P1
0	1	P2

Tagged TLBs

- TLB entries are “tagged” with a *process context identifier*
- Additional bits in CR3 register tells hardware about the current PCID (only available in x86-64)

TLB

Virtual page number	Physical page number	Process context identifier
0	0	P1
2	4	P1
0	1	P2

Tagged TLBs

- TLB entries are “tagged” with a *process context identifier*
- Additional bits in CR3 register tells hardware about the current PCID (only available in x86-64)
- Upon context switch:

TLB

Virtual page number	Physical page number	Process context identifier
0	0	P1
2	4	P1
0	1	P2

Tagged TLBs

- TLB entries are “tagged” with a *process context identifier*
- Additional bits in CR3 register tells hardware about the current PCID (only available in x86-64)
- Upon context switch:
 - OS changes CR3: PCID, page directory base

TLB

Virtual page number	Physical page number	Process context identifier
0	0	P1
2	4	P1
0	1	P2

Tagged TLBs

- TLB entries are “tagged” with a *process context identifier*
- Additional bits in CR3 register tells hardware about the current PCID (only available in x86-64)
- Upon context switch:
 - OS changes CR3: PCID, page directory base
 - Hardware *need not* invalidate TLB entries.
~5us

TLB

Virtual page number	Physical page number	Process context identifier
0	0	P1
2	4	P1
0	1	P2

Tagged TLBs

- TLB entries are “tagged” with a *process context identifier*
- Additional bits in CR3 register tells hardware about the current PCID (only available in x86-64)
- Upon context switch:
 - OS changes CR3: PCID, page directory base
 - Hardware *need not* invalidate TLB entries.
~5us
 - When process gets back control, some of its TLB entries might still be present!

TLB




Virtual page number	Physical page number	Process context identifier
0	0	P1
2	4	P1
0	1	P2

Demand Paging

OSTEP Ch 21-22

Demand paging

Providing illusion of large virtual address spaces






Swap space

Demand paging

Providing illusion of large virtual address spaces

- Swap out a page to disk when physical memory is about to run out of space:






Swap space

Demand paging

Providing illusion of large virtual address spaces

- Swap out a page to disk when physical memory is about to run out of space:
- Mark page as not present in page table






Swap space

Demand paging

Providing illusion of large virtual address spaces

- Swap out a page to disk when physical memory is about to run out of space:
- Mark page as not present in page table
- Remember where page is swapped out on disk





Swap space

Demand paging

Providing illusion of large virtual address spaces

- Swap out a page to disk when physical memory is about to run out of space:
 - Mark page as not present in page table
 - Remember where page is swapped out on disk
- Add swapped out physical page to free list




Swap space

Demand paging

Providing illusion of large virtual address spaces


- Swap out a page to disk when physical memory is about to run out of space:
 - Mark page as not present in page table
 - Remember where page is swapped out on disk
- Add swapped out physical page to free list



Demand paging

Providing illusion of large virtual address spaces


- Swap out a page to disk when physical memory is about to run out of space:
 - Mark page as not present in page table
 - Remember where page is swapped out on disk
- Add swapped out physical page to free list



Demand paging


Providing illusion of large virtual address spaces

- Swap out a page to disk when physical memory is about to run out of space:
 - Mark page as not present in page table
 - Remember where page is swapped out on disk
- Add swapped out physical page to free list



Disk layout with swap space

- Reserved swap blocks, not touched by file system
- Swap space does not require crash consistency: anyways garbage after restart (all processes are dead)



Swapping out a page

- Find a free swap block on disk




Figure 21.1: Physical Memory and Swap Space

Page table entry for Proc 1 [VPN2]

31												12 11 9 8 7 6 5 4 3 2 1 0									
PFN=1 Page Base Address												Avail	G	P A T	D	A	P C D	P W T	U / S	R / W	1

Swapping out a page

- Find a free swap block on disk
- Mark not present, remember swap block number in PTE




Figure 21.1: Physical Memory and Swap Space

Page table entry for Proc 1 [VPN2]

31												12 11 9 8 7 6 5 4 3 2 1 0									
PFN=1 Page Base Address												Avail	G	P A T	D	A	P C D	P W T	U / S	R / W	1

Swapping out a page

- Find a free swap block on disk
- Mark not present, remember swap block number in PTE




Figure 21.1: Physical Memory and Swap Space

Page table entry for Proc 1 [VPN2]

31												12 11 9 8 7 6 5 4 3 2 1 0										
PFN=1 Page Base Address												Avail		G	P A T	D	A	P C D	P W T	U / S	R / W	1

Swapping out a page

- Find a free swap block on disk
- Mark not present, remember swap block number in PTE
- Run INVLPG instruction to remove page from TLB




Figure 21.1: Physical Memory and Swap Space

Page table entry for Proc 1 [VPN2]

31												12 11 9 8 7 6 5 4 3 2 1 0									
PFN=1 Page Base Address												Avail	G	P A T	D	A	P C D	P W T	U / S	R / W	1

Swapping out a page

- Find a free swap block on disk
- Mark not present, remember swap block number in PTE
- Run INVLPG instruction to remove page from TLB
- Copy page to the free block




Figure 21.1: Physical Memory and Swap Space

Page table entry for Proc 1 [VPN2]

31												12		11	9		8	7	6	5	4	3	2	1	0
PFN=1 Page Base Address												Avail		G	P A T	D	A	P C D	P W T	U / S	R / W	1			

Swapping out a page

- Find a free swap block on disk
- Mark not present, remember swap block number in PTE
- Run INVLPG instruction to remove page from TLB
- Copy page to the free block
- Add page to free list




Figure 21.1: Physical Memory and Swap Space

Page table entry for Proc 1 [VPN2]

31												12 11				9	8	7	6	5	4	3	2	1	0
PFN=1												Page Base Address				Avail	G	P A T	D	A	P C D	P W T	U / S	R / W	1

Swapping out a page

- Find a free swap block on disk
- Mark not present, remember swap block number in PTE
- Run INVLPG instruction to remove page from TLB
- Copy page to the free block
- Add page to free list




Figure 21.1: Physical Memory and Swap Space

Page table entry for Proc 1 [VPN2]

31												12		11	9	8	7	6	5	4	3	2	1	0
Page Base Address												Avail		G	P A T	D	A	P C D	P W T	U / S	R / W	P		

Swapping out a page

- Find a free swap block on disk
- Mark not present, remember swap block number in PTE
- Run INVLPG instruction to remove page from TLB
- Copy page to the free block
- Add page to free list




Figure 21.1: Physical Memory and Swap Space

Page table entry for Proc 1 [VPN2]

31												12 11 9 8 7 6 5 4 3 2 1 0									
Swap block # 2 ^{ress}												Avail	G	P A T	D	A	P C D	P W T	U / S	R / W	0

Swapping out a page

- Find a free swap block on disk
- Mark not present, remember swap block number in PTE
- Run INVLPG instruction to remove page from TLB
- Copy page to the free block
- Add page to free list




Figure 21.1: Physical Memory and Swap Space

Page table entry for Proc 1 [VPN2]

31												12 11				9	8	7	6	5	4	3	2	1	0
Swap block # 2 _{ess}												Avail				G		P A T	D	A	P C D	P W T	U / S	R / W	0

Swapping in a page

- Hardware does page table walk to find that the page is not present. It raise page fault with linear address in CR2.




Figure 21.1: Physical Memory and Swap Space

Page table entry for Proc 1 [VPN2]



Swapping in a page

- Hardware does page table walk to find that the page is not present. It raise page fault with linear address in CR2.
- OS handles page fault:




Figure 21.1: Physical Memory and Swap Space

Page table entry for Proc 1 [VPN2]

31																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																											
----	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--	--

Swapping in a page

- Hardware does page table walk to find that the page is not present. It raise page fault with linear address in CR2.
- OS handles page fault:
 - Copies page to a free physical page




Figure 21.1: Physical Memory and Swap Space

Page table entry for Proc 1 [VPN2]



Swapping in a page

- Hardware does page table walk to find that the page is not present. It raise page fault with linear address in CR2.
- OS handles page fault:
 - Copies page to a free physical page
 - Updates page table entry





Figure 21.1: Physical Memory and Swap Space

Page table entry for Proc 1 [VPN2]



Swapping in a page

- Hardware does page table walk to find that the page is not present. It raise page fault with linear address in CR2.
- OS handles page fault:
 - Copies page to a free physical page
 - Updates page table entry




Figure 21.1: Physical Memory and Swap Space

Page table entry for Proc 1 [VPN2]



Swapping in a page

- Hardware does page table walk to find that the page is not present. It raise page fault with linear address in CR2.
- OS handles page fault:
 - Copies page to a free physical page
 - Updates page table entry
- Hardware retries instruction. This time finds the page, adds entry to TLB, continues as normal




Figure 21.1: Physical Memory and Swap Space

Page table entry for Proc 1 [VPN2]



Page replacement policies

- When to evict pages?
- How many pages to evict?
- Which page to evict?

When to evict pages? How many pages to evict?

- Swap out one page when we completely run out of physical memory
 - What if OS itself needed a new page?

When to evict pages? How many pages to evict?

- Swap out one page when we completely run out of physical memory
 - What if OS itself needed a new page?
- Start swapping out before we completely run out
 - When there are less than N free pages left

When to evict pages? How many pages to evict?

- Swap out one page when we completely run out of physical memory
 - What if OS itself needed a new page?
- Start swapping out before we completely run out
 - When there are less than N free pages left
- Swap out multiple pages in one shot until we have $M (> N)$ free pages left
 - Sends multiple disk writes in one shot reduces seek delay

Which page to evict?

- Goal: minimize number of swap ins/outs

Page access sequence:

0, 1, 2, 0, 1, 3, 0, 3, 1, 2, 1

Which page to evict?

- Goal: minimize number of swap ins/outs
- Belady's algorithm for optimal page replacement
 - Evict the page required furthest in the future
 - Optimal because all other pages will be required sooner

Page access sequence:

0, 1, 2, 0, 1, 3, 0, 3, 1, 2, 1

Which page to evict?

- Goal: minimize number of swap ins/outs
- Belady's algorithm for optimal page replacement
 - Evict the page required furthest in the future
 - Optimal because all other pages will be required sooner

Page access sequence:

0, 1, 2, 0, 1, 3, 0, 3, 1, 2, 1

Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		0

Which page to evict?

- Goal: minimize number of swap ins/outs
- Belady's algorithm for optimal page replacement
 - Evict the page required furthest in the future
 - Optimal because all other pages will be required sooner

Page access sequence:

0, 1, 2, 0, 1, 3, 0, 3, 1, 2, 1

Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		0
1	Miss		0, 1
2	Miss		0, 1, 2
0	Hit		0, 1, 2
1	Hit		0, 1, 2

Which page to evict?

- Goal: minimize number of swap ins/outs
- Belady's algorithm for optimal page replacement
 - Evict the page required furthest in the future
 - Optimal because all other pages will be required sooner

Page access sequence:

0, 1, 2, 0, 1, 3, 0, 3, 1, 2, 1

Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		0
1	Miss		0, 1
2	Miss		0, 1, 2
0	Hit		0, 1, 2
1	Hit		0, 1, 2
3	Miss	2	0, 1, 3
0	Hit		0, 1, 3
3	Hit		0, 1, 3
1	Hit		0, 1, 3
2	Miss		0, 1, 3
1	Hit		0, 1, 3

Which page to evict?

- Goal: minimize number of swap ins/outs
- Belady's algorithm for optimal page replacement
 - Evict the page required furthest in the future
 - Optimal because all other pages will be required sooner

Page access sequence:

0, 1, 2, 0, 1, 3, 0, 3, 1, 2, 1

Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		0
1	Miss		0, 1
2	Miss		0, 1, 2
0	Hit		0, 1, 2
1	Hit		0, 1, 2
3	Miss	2	0, 1, 3
0	Hit		0, 1, 3
3	Hit		0, 1, 3
1	Hit		0, 1, 3
2	Miss	3	0, 1, 2

Which page to evict?

- Goal: minimize number of swap ins/outs
- Belady's algorithm for optimal page replacement
 - Evict the page required furthest in the future
 - Optimal because all other pages will be required sooner

Page access sequence:

0, 1, 2, 0, 1, 3, 0, 3, 1, 2, 1

Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		0
1	Miss		0, 1
2	Miss		0, 1, 2
0	Hit		0, 1, 2
1	Hit		0, 1, 2
3	Miss	2	0, 1, 3
0	Hit		0, 1, 3
3	Hit		0, 1, 3
1	Hit		0, 1, 3
2	Miss	3	0, 1, 2
1	Hit		0, 1, 2

Figure 22.1: Tracing The Optimal Policy

Which page to evict?

- Goal: minimize number of swap ins/outs
- Belady's algorithm for optimal page replacement
 - Evict the page required furthest in the future
 - Optimal because all other pages will be required sooner
- Future is unknown!

Page access sequence:

0, 1, 2, 0, 1, 3, 0, 3, 1, 2, 1

Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		0
1	Miss		0, 1
2	Miss		0, 1, 2
0	Hit		0, 1, 2
1	Hit		0, 1, 2
3	Miss	2	0, 1, 3
0	Hit		0, 1, 3
3	Hit		0, 1, 3
1	Hit		0, 1, 3
2	Miss	3	0, 1, 2
1	Hit		0, 1, 2

Figure 22.1: Tracing The Optimal Policy

FIFO

- Evict the page that came first to the cache
- OS appends the page to a queue when it swaps in a page (or when it allocates a new page)

Memory access sequence:

0, 1, 2, 0, 1, 3, 0, 3, 1, 2, 1

Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		First-in→ 0

FIFO

- Evict the page that came first to the cache
- OS appends the page to a queue when it swaps in a page (or when it allocates a new page)

Memory access sequence:

0, 1, 2, 0, 1, 3, 0, 3, 1, 2, 1

Access	Hit/Miss?	Evict	Resulting Cache State	
0	Miss		First-in→	0
1	Miss		First-in→	0, 1
2	Miss		First-in→	0, 1, 2
0	Hit		First-in→	0, 1, 2
1	Hit		First-in→	0, 1, 2

FIFO

- Evict the page that came first to the cache
- OS appends the page to a queue when it swaps in a page (or when it allocates a new page)

Memory access sequence:

0, 1, 2, 0, 1, 3, 0, 3, 1, 2, 1

Access	Hit/Miss?	Evict	Resulting Cache State	
0	Miss		First-in→	0
1	Miss		First-in→	0, 1
2	Miss		First-in→	0, 1, 2
0	Hit		First-in→	0, 1, 2
1	Hit		First-in→	0, 1, 2
3	Miss	0	First-in→	1, 2, 3

FIFO

- Evict the page that came first to the cache
- OS appends the page to a queue when it swaps in a page (or when it allocates a new page)

Memory access sequence:

0, 1, 2, 0, 1, 3, 0, 3, 1, 2, 1

Access	Hit/Miss?	Evict	Resulting Cache State	
0	Miss		First-in→	0
1	Miss		First-in→	0, 1
2	Miss		First-in→	0, 1, 2
0	Hit		First-in→	0, 1, 2
1	Hit		First-in→	0, 1, 2
3	Miss	0	First-in→	1, 2, 3
0	Miss	1	First-in→	2, 3, 0

FIFO

- Evict the page that came first to the cache
- OS appends the page to a queue when it swaps in a page (or when it allocates a new page)

Memory access sequence:

0, 1, 2, 0, 1, 3, 0, 3, 1, 2, 1

Access	Hit/Miss?	Evict	Resulting Cache State	
0	Miss		First-in→	0
1	Miss		First-in→	0, 1
2	Miss		First-in→	0, 1, 2
0	Hit		First-in→	0, 1, 2
1	Hit		First-in→	0, 1, 2
3	Miss	0	First-in→	1, 2, 3
0	Miss	1	First-in→	2, 3, 0
3	Hit		First-in→	2, 3, 0

FIFO

- Evict the page that came first to the cache
- OS appends the page to a queue when it swaps in a page (or when it allocates a new page)

Memory access sequence:

0, 1, 2, 0, 1, 3, 0, 3, 1, 2, 1

Access	Hit/Miss?	Evict	Resulting Cache State	
0	Miss		First-in→	0
1	Miss		First-in→	0, 1
2	Miss		First-in→	0, 1, 2
0	Hit		First-in→	0, 1, 2
1	Hit		First-in→	0, 1, 2
3	Miss	0	First-in→	1, 2, 3
0	Miss	1	First-in→	2, 3, 0
3	Hit		First-in→	2, 3, 0
1	Miss	2	First-in→	3, 0, 1
2	Miss	3	First-in→	0, 1, 2
1	Hit		First-in→	0, 1, 2

Figure 22.2: Tracing The FIFO Policy

Belady's anomaly

Bigger caches can have lower hit rates!

Access	Hit/miss	Resulting cache	Access	Hit/miss	Resulting cache
1	Miss	1	1	Miss	1
2	Miss	1, 2	2	Miss	1, 2
3	Miss	1, 2, 3	3	Miss	1, 2, 3
4	Miss	2, 3, 4	4	Miss	1, 2, 3, 4
1	Miss	3, 4, 1	1	Hit	1, 2, 3, 4
2	Miss	4, 1, 2	2	Hit	1, 2, 3, 4
5	Miss	1, 2, 5	5	Miss	2, 3, 4, 5
1	Hit	1, 2, 5	1	Miss	3, 4, 5, 1
2	Hit	1, 2, 5	2	Miss	4, 5, 1, 2
3	Miss	2, 5, 3	3	Miss	5, 1, 2, 3
4	Miss	5, 3, 4	4	Miss	1, 2, 3, 4
5	Hit	5, 3, 4	5	Miss	2, 3, 4, 5

Belady's anomaly

Bigger caches can have lower hit rates!

Access	Hit/miss	Resulting cache	Access	Hit/miss	Resulting cache
1	Miss	1	1	Miss	1
2	Miss	1, 2	2	Miss	1, 2
3	Miss	1, 2, 3	3	Miss	1, 2, 3
4	Miss	2, 3, 4	4	Miss	1, 2, 3, 4
1	Miss	3, 4, 1	1	Hit	1, 2, 3, 4
2	Miss	4, 1, 2	2	Hit	1, 2, 3, 4
5	Miss	1, 2, 5	5	Miss	2, 3, 4, 5
1	Hit	1, 2, 5	1	Miss	3, 4, 5, 1
2	Hit	1, 2, 5	2	Miss	4, 5, 1, 2
3	Miss	2, 5, 3	3	Miss	5, 1, 2, 3
4	Miss	5, 3, 4	4	Miss	1, 2, 3, 4
5	Hit	5, 3, 4	5	Miss	2, 3, 4, 5

FIFO does not follow “stack property”. Cache of size 4 may not contain elements in cache of size 3.

Belady's anomaly

Bigger caches can have lower hit rates!

Access	Hit/miss	Resulting cache	Access	Hit/miss	Resulting cache
1	Miss	1	1	Miss	1
2	Miss	1, 2	2	Miss	1, 2
3	Miss	1, 2, 3	3	Miss	1, 2, 3
4	Miss	2, 3, 4	4	Miss	1, 2, 3, 4
1	Miss	3, 4, 1	1	Hit	1, 2, 3, 4
2	Miss	4, 1, 2	2	Hit	1, 2, 3, 4
5	Miss	1, 2, 5	5	Miss	2, 3, 4, 5
1	Hit	1, 2, 5	1	Miss	3, 4, 5, 1
2	Hit	1, 2, 5	2	Miss	4, 5, 1, 2
3	Miss	2, 5, 3	3	Miss	5, 1, 2, 3
4	Miss	5, 3, 4	4	Miss	1, 2, 3, 4
5	Hit	5, 3, 4	5	Miss	2, 3, 4, 5

FIFO does not follow “stack property”. Cache of size 4 may not contain elements in cache of size 3.

Fairness in page replacement

- Someone had lots of pages. I had very little. My page was evicted



Fairness in page replacement

- Someone had lots of pages. I had very little. My page was evicted
- OS maintains “resident size” per process: 1, 7, 9



Fairness in page replacement

- Someone had **lots of pages**. I had **very little**. My page was evicted
- OS maintains “resident size” per process: 1, 7, 9
- First select a **victim process** with highest resident size, remove its pages



Least Recently Used (LRU)

- Most programs exhibit temporal locality:
 - If a page was accessed recently, it shall be accessed soon
 - Keep list according to recency

Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		LRU→ 0

Least Recently Used (LRU)

- Most programs exhibit temporal locality:
 - If a page was accessed recently, it shall be accessed soon
- Keep list according to recency

Access	Hit/Miss?	Evict	Resulting Cache State	
0	Miss		LRU→	0
1	Miss		LRU→	0, 1
2	Miss		LRU→	0, 1, 2
0	Hit		LRU→	1, 2, 0
1	Hit		LRU→	2, 0, 1

Least Recently Used (LRU)

- Most programs exhibit temporal locality:
 - If a page was accessed recently, it shall be accessed soon
- Keep list according to recency

Access	Hit/Miss?	Evict	Resulting Cache State	
0	Miss		LRU→	0
1	Miss		LRU→	0, 1
2	Miss		LRU→	0, 1, 2
0	Hit		LRU→	1, 2, 0
1	Hit		LRU→	2, 0, 1
3	Miss	2	LRU→	0, 1, 3

Least Recently Used (LRU)

- Most programs exhibit temporal locality:
 - If a page was accessed recently, it shall be accessed soon
- Keep list according to recency

Access	Hit/Miss?	Evict	Resulting Cache State	
0	Miss		LRU→	0
1	Miss		LRU→	0, 1
2	Miss		LRU→	0, 1, 2
0	Hit		LRU→	1, 2, 0
1	Hit		LRU→	2, 0, 1
3	Miss	2	LRU→	0, 1, 3
0	Hit		LRU→	1, 3, 0
3	Hit		LRU→	1, 0, 3
1	Hit		LRU→	0, 3, 1

Least Recently Used (LRU)

- Most programs exhibit temporal locality:
 - If a page was accessed recently, it shall be accessed soon
- Keep list according to recency

Access	Hit/Miss?	Evict	Resulting Cache State	
0	Miss		LRU→	0
1	Miss		LRU→	0, 1
2	Miss		LRU→	0, 1, 2
0	Hit		LRU→	1, 2, 0
1	Hit		LRU→	2, 0, 1
3	Miss	2	LRU→	0, 1, 3
0	Hit		LRU→	1, 3, 0
3	Hit		LRU→	1, 0, 3
1	Hit		LRU→	0, 3, 1
2	Miss	0	LRU→	3, 1, 2

Least Recently Used (LRU)

- Most programs exhibit temporal locality:
 - If a page was accessed recently, it shall be accessed soon
- Keep list according to recency

Access	Hit/Miss?	Evict	Resulting Cache State	
0	Miss		LRU→	0
1	Miss		LRU→	0, 1
2	Miss		LRU→	0, 1, 2
0	Hit		LRU→	1, 2, 0
1	Hit		LRU→	2, 0, 1
3	Miss	2	LRU→	0, 1, 3
0	Hit		LRU→	1, 3, 0
3	Hit		LRU→	1, 0, 3
1	Hit		LRU→	0, 3, 1
2	Miss	0	LRU→	3, 1, 2
1	Hit		LRU→	3, 2, 1

Figure 22.5: Tracing The LRU Policy

Difficulty in implementing LRU

- In FIFO, list is updated by the OS when a new page is allocated, or when a page is swapped out

Access	Hit/Miss?	Evict	Resulting Cache State	
0	Miss		LRU→	0
1	Miss		LRU→	0, 1
2	Miss		LRU→	0, 1, 2
0	Hit		LRU→	1, 2, 0
1	Hit		LRU→	2, 0, 1
3	Miss	2	LRU→	0, 1, 3
0	Hit		LRU→	1, 3, 0
3	Hit		LRU→	1, 0, 3
1	Hit		LRU→	0, 3, 1
2	Miss	0	LRU→	3, 1, 2
1	Hit		LRU→	3, 2, 1

Figure 22.5: Tracing The LRU Policy

Difficulty in implementing LRU

- In FIFO, list is updated by the OS when a new page is allocated, or when a page is swapped out
- In LRU, list needs to be updated at every access
- OS is not running during page accesses :-/

Access	Hit/Miss?	Evict	Resulting Cache State	
0	Miss		LRU→	0
1	Miss		LRU→	0, 1
2	Miss		LRU→	0, 1, 2
0	Hit		LRU→	1, 2, 0
1	Hit		LRU→	2, 0, 1
3	Miss	2	LRU→	0, 1, 3
0	Hit		LRU→	1, 3, 0
3	Hit		LRU→	1, 0, 3
1	Hit		LRU→	0, 3, 1
2	Miss	0	LRU→	3, 1, 2
1	Hit		LRU→	3, 2, 1

Figure 22.5: Tracing The LRU Policy

Implementation options

- Option 1: Hardware maintains the LRU list

Implementation options

- Option 1: Hardware maintains the LRU list
 - 4GB / 4KB $\sim 2^{20}$ pages

Implementation options

- Option 1: Hardware maintains the LRU list
 - 4GB / 4KB $\sim 2^{20}$ pages
 - List size: 20 bits * 2^{20} pages $\sim 3\text{MB}$

Implementation options

- Option 1: Hardware maintains the LRU list
 - 4GB / 4KB $\sim 2^{20}$ pages
 - List size: 20 bits * 2^{20} pages $\sim 3\text{MB}$
 - List cannot be in CPU $\Rightarrow$ must be in memory

Implementation options

- Option 1: Hardware maintains the LRU list
 - 4GB / 4KB $\sim 2^{20}$ pages
 - List size: 20 bits * 2^{20} pages $\sim 3\text{MB}$
 - List cannot be in CPU $\Rightarrow$ must be in memory
 - Each memory access causes another set of memory accesses to update list

Implementation options

- Option 2: Hardware updates timestamp in page table entries.
OS scans PTEs to find page with oldest timestamp

Implementation options

- Option 2: Hardware updates timestamp in page table entries.
OS scans PTEs to find page with oldest timestamp

Page table

Physical page number	Access Timestamp
10	8:11am
11	7:05am
15	7:00am
12	8:21am

Implementation options

- Option 2: Hardware updates timestamp in page table entries. OS scans PTEs to find page with oldest timestamp
- PTEs live in memory. Updating timestamp again touches memory at every access. Defeats TLB.

Page table

Physical page number	Access Timestamp
10	8:11am
11	7:05am
15	7:00am
12	8:21am

Implementation options

- Option 2: Hardware updates timestamp in page table entries. OS scans PTEs to find page with oldest timestamp
- PTEs live in memory. Updating timestamp again touches memory at every access. Defeats TLB.

Page table

Physical page number	Access Timestamp
10	8:11am
11	7:05am
15	7:00am
12	8:21am

TLB

Virtual page number	Physical page number
0	10
1	11
2	15
3	12

Implementation options

- Option 2: Hardware updates timestamp in page table entries. OS scans PTEs to find page with oldest timestamp
- PTEs live in memory. Updating timestamp again touches memory at every access. Defeats TLB.
- Lazily update timestamp:

Page table

Physical page number	Access Timestamp
10	8:11am
11	7:05am
15	7:00am
12	8:21am

TLB

Virtual page number	Physical page number
0	10
1	11
2	15
3	12

Implementation options

- Option 2: Hardware updates timestamp in page table entries. OS scans PTEs to find page with oldest timestamp
- PTEs live in memory. Updating timestamp again touches memory at every access. Defeats TLB.
- Lazily update timestamp:
 - when mapping is brought to TLB

Page table

Physical page number	Access Timestamp
10	8:11am
11	7:05am
15	7:00am
12	8:21am

TLB

Virtual page number	Physical page number
0	10
1	11
2	15
3	12

Implementation options

- Option 2: Hardware updates timestamp in page table entries. OS scans PTEs to find page with oldest timestamp
- PTEs live in memory. Updating timestamp again touches memory at every access. Defeats TLB.
- Lazily update timestamp:
 - when mapping is brought to TLB
 - when mapping is evicted from TLB

Page table

Physical page number	Access Timestamp
10	8:11am
11	7:05am
15	7:00am
12	8:21am

TLB

Virtual page number	Physical page number
0	10
1	11
2	15
3	12

Implementation options

- Option 2: Hardware updates timestamp in page table entries. OS scans PTEs to find page with oldest timestamp
- PTEs live in memory. Updating timestamp again touches memory at every access. Defeats TLB.
- Lazily update timestamp:
 - when mapping is brought to TLB
 - when mapping is evicted from TLB
 - Once in a while

Page table

Physical page number	Access Timestamp
10	8:11am
11	7:05am
15	7:00am
12	8:21am

TLB

Virtual page number	Physical page number
0	10
1	11
2	15
3	12

Implementation options

- Option 2: Hardware updates timestamp in page table entries. OS scans PTEs to find page with oldest timestamp
- PTEs live in memory. Updating timestamp again touches memory at every access. Defeats TLB.
- Lazily update timestamp:
 - when mapping is brought to TLB
 - when mapping is evicted from TLB
 - Once in a while
- Need many more bits in PTE to store timestamp

Page table

Physical page number	Access Timestamp
10	8:11am
11	7:05am
15	7:00am
12	8:21am

TLB

Virtual page number	Physical page number
0	10
1	11
2	15
3	12

Implementation options

- Option 2: Hardware updates timestamp in page table entries. OS scans PTEs to find page with oldest timestamp
- PTEs live in memory. Updating timestamp again touches memory at every access. Defeats TLB.
- Lazily update timestamp:
 - when mapping is brought to TLB
 - when mapping is evicted from TLB
 - Once in a while
- Need many more bits in PTE to store timestamp
- Victim process has highest resident size. Scanning (worst case 2^{20}) PTEs will be slow

Page table

Physical page number	Access Timestamp
10	8:11am
11	7:05am
15	7:00am
12	8:21am

TLB

Virtual page number	Physical page number
0	10
1	11
2	15
3	12

Approximating LRU

- Give up on finding *least* recently used. OK to evict a *less* recently used

Page table

Physical page number	Access Timestamp
10	8:11am
11	7:05am
15	7:00am
12	8:21am

Approximating LRU

- Give up on finding *least* recently used. OK to evict a *less* recently used

Page table

Physical page number	Access Timestamp
10	8:11am
11	7:05am
15	7:00am
12	8:21am

Page-Table Entry (4-KByte Page)




Figure 3-14. Format of Page-Directory and Page-Table Entries for 4-KByte Pages and 32-Bit Physical Addresses

Approximating LRU

- Give up on finding *least* recently used. OK to evict a *less* recently used
- Hardware just lazily sets 1 access bit

Page table

Physical page number	Access Timestamp
10	8:11am
11	7:05am
15	7:00am
12	8:21am




Figure 3-14. Format of Page-Directory and Page-Table Entries for 4-KByte Pages and 32-Bit Physical Addresses

Clock algorithm

- OS clears access bit. Evicts page with access bit = 0
- Hardware sets access bit to 1
- Evicted page was “not recently used”

Victim proc's pages

	AB
	1
	1
	1
	1
OS →	1
	1
	1
	1

Clock algorithm

- OS clears access bit. Evicts page with access bit = 0
- Hardware sets access bit to 1
- Evicted page was “not recently used”

Victim proc's pages

	AB
	1
	1
	1
	1
OS →	0
	1
	1
	1

Clock algorithm

- OS clears access bit. Evicts page with access bit = 0
- Hardware sets access bit to 1
- Evicted page was “not recently used”

Victim proc's pages

	AB
	1
	1
	1
	1
	0
OS →	1
	1
	1

Clock algorithm

- OS clears access bit. Evicts page with access bit = 0
- Hardware sets access bit to 1
- Evicted page was “not recently used”

Victim proc's pages

	AB
	1
	1
	1
	1
	0
OS →	0
	1
	1

Clock algorithm

- OS clears access bit. Evicts page with access bit = 0
- Hardware sets access bit to 1
- Evicted page was “not recently used”

Victim proc's pages

	AB
	1
	1
	1
	1
	0
	0
OS →	1
	1

Clock algorithm

- OS clears access bit. Evicts page with access bit = 0
- Hardware sets access bit to 1
- Evicted page was “not recently used”

Victim proc's pages

	AB
	1
	1
	1
	1
	0
	0
OS →	0
	1

Clock algorithm

- OS clears access bit. Evicts page with access bit = 0
- Hardware sets access bit to 1
- Evicted page was “not recently used”

Victim proc's pages	
	AB
	1
	1
	1
	1
	0
	0
	0
OS →	1

Clock algorithm

- OS clears access bit. Evicts page with access bit = 0
- Hardware sets access bit to 1
- Evicted page was “not recently used”

Victim proc's pages	
	AB
	1
	1
	1
	1
	0
	0
	0
OS →	0

Clock algorithm

- OS clears access bit. Evicts page with access bit = 0
- Hardware sets access bit to 1
- Evicted page was “not recently used”

OS →

	AB
	1
	1
	1
	1
	0
	0
	0
	0

Clock algorithm

- OS clears access bit. Evicts page with access bit = 0
- Hardware sets access bit to 1
- Evicted page was “not recently used”

Victim proc's pages

	AB
OS →	1
	1
	1
	1
	0
	0
	1
	0

Clock algorithm

- OS clears access bit. Evicts page with access bit = 0
- Hardware sets access bit to 1
- Evicted page was “not recently used”

Victim proc's pages

	AB
OS →	0
	1
	1
	1
	0
	0
	1
	0

Clock algorithm

- OS clears access bit. Evicts page with access bit = 0
- Hardware sets access bit to 1
- Evicted page was “not recently used”

Victim proc's pages

	AB
	0
OS →	1
	1
	1
	0
	0
	1
	0

Clock algorithm

- OS clears access bit. Evicts page with access bit = 0
- Hardware sets access bit to 1
- Evicted page was “not recently used”

Victim proc's pages

	AB
	0
OS →	0
	1
	1
	0
	0
	1
	0

Clock algorithm

- OS clears access bit. Evicts page with access bit = 0
- Hardware sets access bit to 1
- Evicted page was “not recently used”

Victim proc's pages

	AB
	0
OS →	0
	1
	1
	1
	0
	1
	0

Clock algorithm

- OS clears access bit. Evicts page with access bit = 0
- Hardware sets access bit to 1
- Evicted page was “not recently used”

Victim proc's pages

	AB
	0
	0
OS →	1
	1
	1
	0
	1
	0

Clock algorithm

- OS clears access bit. Evicts page with access bit = 0
- Hardware sets access bit to 1
- Evicted page was “not recently used”

Victim proc's pages

	AB
	0
	0
OS →	0
	1
	1
	0
	1
	0

Clock algorithm

- OS clears access bit. Evicts page with access bit = 0
- Hardware sets access bit to 1
- Evicted page was “not recently used”

Victim proc's pages

	AB
	0
	0
	0
OS →	1
	1
	0
	1
	0

Clock algorithm

- OS clears access bit. Evicts page with access bit = 0
- Hardware sets access bit to 1
- Evicted page was “not recently used”

Victim proc's pages

	AB
	0
	0
	0
OS →	0
	1
	0
	1
	0

Clock algorithm

- OS clears access bit. Evicts page with access bit = 0
- Hardware sets access bit to 1
- Evicted page was “not recently used”

Victim proc's pages

	AB
	0
	0
	0
	0
OS →	1
	0
	1
	0

Clock algorithm

- OS clears access bit. Evicts page with access bit = 0
- Hardware sets access bit to 1
- Evicted page was “not recently used”

Victim proc's pages

	AB
	0
	0
	0
	0
OS →	0
	0
	1
	0

Clock algorithm

- OS clears access bit. Evicts page with access bit = 0
- Hardware sets access bit to 1
- Evicted page was “not recently used”

Victim proc's pages

	AB
	0
	0
	0
	0
	0
OS →	0
	1
	0

Clock algorithm

- OS clears access bit. Evicts page with access bit = 0
- Hardware sets access bit to 1
- Evicted page was “not recently used”

Victim proc's pages

	AB
	0
	0
	0
	0
	0
OS →	0
	1
	0

Clock algorithm (2)

- Optimisation:
 - We should prefer evicting page that has not changed since it was brought back from disk

Victim proc's pages

	DB	AB
	0	1
	0	1
	0	1
	1	1
OS →	0	1
	1	1
	1	1
	0	1

Clock algorithm (2)

- Optimisation:
 - We should prefer evicting page that has not changed since it was brought back from disk
 - Such pages can just be deleted without doing a copy

Victim proc's pages

	DB	AB
	0	1
	0	1
	0	1
	1	1
OS →	0	1
	1	1
	1	1
	0	1

Clock algorithm (2)

- Optimisation:
 - We should prefer evicting page that has not changed since it was brought back from disk
 - Such pages can just be deleted without doing a copy
- OS clears dirty bit when it brings a page to memory from disk

Victim proc's pages

	DB	AB
	0	1
	0	1
	0	1
	1	1
OS →	0	1
	1	1
	1	1
	0	1

Clock algorithm (2)

- Optimisation:
 - We should prefer evicting page that has not changed since it was brought back from disk
 - Such pages can just be deleted without doing a copy
- OS clears dirty bit when it brings a page to memory from disk
- Evict dirty pages iff not able to find a clean page

Victim proc's pages

	DB	AB
	0	1
	0	1
	0	1
	1	1
OS →	0	1
	1	1
	1	1
	0	1

Clock algorithm (2)

- Optimisation:
 - We should prefer evicting page that has not changed since it was brought back from disk
 - Such pages can just be deleted without doing a copy
- OS clears dirty bit when it brings a page to memory from disk
- Evict dirty pages iff not able to find a clean page

Victim proc's pages

	DB	AB
	0	1
	0	1
	0	1
	1	1
OS →	0	0
	1	1
	1	1
	0	1

Clock algorithm (2)

- Optimisation:
 - We should prefer evicting page that has not changed since it was brought back from disk
 - Such pages can just be deleted without doing a copy
- OS clears dirty bit when it brings a page to memory from disk
- Evict dirty pages iff not able to find a clean page

Victim proc's pages

	DB	AB
	0	1
	0	1
	0	1
	1	1
	0	0
OS →	1	1
	1	1
	0	1

Clock algorithm (2)

- Optimisation:
 - We should prefer evicting page that has not changed since it was brought back from disk
 - Such pages can just be deleted without doing a copy
- OS clears dirty bit when it brings a page to memory from disk
- Evict dirty pages iff not able to find a clean page

Victim proc's pages

	DB	AB
	0	1
	0	1
	0	1
	1	1
	0	0
OS →	1	0
	1	1
	0	1

Clock algorithm (2)

- Optimisation:
 - We should prefer evicting page that has not changed since it was brought back from disk
 - Such pages can just be deleted without doing a copy
- OS clears dirty bit when it brings a page to memory from disk
- Evict dirty pages iff not able to find a clean page

Victim proc's pages

	DB	AB
	0	1
	0	1
	0	1
	1	1
	0	0
	1	0
OS →	1	1
	0	1

Clock algorithm (2)

- Optimisation:
 - We should prefer evicting page that has not changed since it was brought back from disk
 - Such pages can just be deleted without doing a copy
- OS clears dirty bit when it brings a page to memory from disk
- Evict dirty pages iff not able to find a clean page

Victim proc's pages

	DB	AB
	0	1
	0	1
	0	1
	1	1
	0	0
	1	0
OS →	1	0
	0	1

Clock algorithm (2)

- Optimisation:
 - We should prefer evicting page that has not changed since it was brought back from disk
 - Such pages can just be deleted without doing a copy
- OS clears dirty bit when it brings a page to memory from disk
- Evict dirty pages iff not able to find a clean page

Victim proc's pages

	DB	AB
	0	1
	0	1
	0	1
	1	1
	0	0
	1	0
	1	0
OS →	0	1

Clock algorithm (2)

- Optimisation:
 - We should prefer evicting page that has not changed since it was brought back from disk
 - Such pages can just be deleted without doing a copy
- OS clears dirty bit when it brings a page to memory from disk
- Evict dirty pages iff not able to find a clean page

Victim proc's pages

	DB	AB
	0	1
	0	1
	0	1
	1	1
	0	0
	1	0
	1	0
OS →	0	0

Clock algorithm (2)

- Optimisation:
 - We should prefer evicting page that has not changed since it was brought back from disk
 - Such pages can just be deleted without doing a copy
- OS clears dirty bit when it brings a page to memory from disk
- Evict dirty pages iff not able to find a clean page

OS →

	DB	AB
	0	1
	0	1
	0	1
	1	1
	0	0
	1	0
	1	0
	0	0

Clock algorithm (2)

- Optimisation:
 - We should prefer evicting page that has not changed since it was brought back from disk
 - Such pages can just be deleted without doing a copy
- OS clears dirty bit when it brings a page to memory from disk
- Evict dirty pages iff not able to find a clean page

OS →

	DB	AB
	0	1
	0	1
	0	1
	1	1
	0	0
	1	0
	1	1
	0	0

Clock algorithm (2)

- Optimisation:
 - We should prefer evicting page that has not changed since it was brought back from disk
 - Such pages can just be deleted without doing a copy
- OS clears dirty bit when it brings a page to memory from disk
- Evict dirty pages iff not able to find a clean page

OS →

	DB	AB
	0	0
	0	1
	0	1
	1	1
	0	0
	1	0
	1	1
	0	0

Clock algorithm (2)

- Optimisation:
 - We should prefer evicting page that has not changed since it was brought back from disk
 - Such pages can just be deleted without doing a copy
- OS clears dirty bit when it brings a page to memory from disk
- Evict dirty pages iff not able to find a clean page

OS →

	DB	AB
	0	0
	0	1
	0	1
	1	1
	0	0
	1	0
	1	1
	0	0

Clock algorithm (2)

- Optimisation:
 - We should prefer evicting page that has not changed since it was brought back from disk
 - Such pages can just be deleted without doing a copy
- OS clears dirty bit when it brings a page to memory from disk
- Evict dirty pages iff not able to find a clean page

Victim proc's pages

	DB	AB
	0	0
OS →	0	0
	0	1
	1	1
	0	0
	1	0
	1	1
	0	0

Clock algorithm (2)

- Optimisation:
 - We should prefer evicting page that has not changed since it was brought back from disk
 - Such pages can just be deleted without doing a copy
- OS clears dirty bit when it brings a page to memory from disk
- Evict dirty pages iff not able to find a clean page

OS →

	DB	AB
	0	0
	0	0
	0	1
	1	1
	0	1
	1	0
	1	1
	0	0

Clock algorithm (2)

- Optimisation:
 - We should prefer evicting page that has not changed since it was brought back from disk
 - Such pages can just be deleted without doing a copy
- OS clears dirty bit when it brings a page to memory from disk
- Evict dirty pages iff not able to find a clean page

Victim proc's pages

	DB	AB
	0	0
	0	0
OS →	0	1
	1	1
	0	1
	1	0
	1	1
	0	0

Clock algorithm (2)

- Optimisation:
 - We should prefer evicting page that has not changed since it was brought back from disk
 - Such pages can just be deleted without doing a copy
- OS clears dirty bit when it brings a page to memory from disk
- Evict dirty pages iff not able to find a clean page

Victim proc's pages

	DB	AB
	0	0
	0	0
OS →	0	0
	1	1
	0	1
	1	0
	1	1
	0	0

Clock algorithm (2)

- Optimisation:
 - We should prefer evicting page that has not changed since it was brought back from disk
 - Such pages can just be deleted without doing a copy
- OS clears dirty bit when it brings a page to memory from disk
- Evict dirty pages iff not able to find a clean page

Victim proc's pages

	DB	AB
	0	0
	0	0
	0	0
OS →	1	1
	0	1
	1	0
	1	1
	0	0

Clock algorithm (2)

- Optimisation:
 - We should prefer evicting page that has not changed since it was brought back from disk
 - Such pages can just be deleted without doing a copy
- OS clears dirty bit when it brings a page to memory from disk
- Evict dirty pages iff not able to find a clean page

Victim proc's pages

	DB	AB
	0	0
	0	0
	0	0
OS →	1	0
	0	1
	1	0
	1	1
	0	0

Clock algorithm (2)

- Optimisation:
 - We should prefer evicting page that has not changed since it was brought back from disk
 - Such pages can just be deleted without doing a copy
- OS clears dirty bit when it brings a page to memory from disk
- Evict dirty pages iff not able to find a clean page

OS →

	DB	AB
	0	0
	0	0
	0	0
	1	0
	0	1
	1	0
	1	1
	0	0

Clock algorithm (2)

- Optimisation:
 - We should prefer evicting page that has not changed since it was brought back from disk
 - Such pages can just be deleted without doing a copy
- OS clears dirty bit when it brings a page to memory from disk
- Evict dirty pages iff not able to find a clean page

Victim proc's pages

	DB	AB
	0	0
	0	0
	0	0
	1	0
OS →	0	0
	1	0
	1	1
	0	0

Clock algorithm (2)

- Optimisation:
 - We should prefer evicting page that has not changed since it was brought back from disk
 - Such pages can just be deleted without doing a copy
- OS clears dirty bit when it brings a page to memory from disk
- Evict dirty pages iff not able to find a clean page

Victim proc's pages

	DB	AB
	0	0
	0	0
	0	0
	1	0
	0	0
OS →	1	0
	1	1
	0	0

Clock algorithm (2)

- Optimisation:
 - We should prefer evicting page that has not changed since it was brought back from disk
 - Such pages can just be deleted without doing a copy
- OS clears dirty bit when it brings a page to memory from disk
- Evict dirty pages iff not able to find a clean page

Victim proc's pages

	DB	AB
	0	0
	0	0
	0	0
	1	0
	0	0
	1	0
OS →	1	1
	0	0

Clock algorithm (2)

- Optimisation:
 - We should prefer evicting page that has not changed since it was brought back from disk
 - Such pages can just be deleted without doing a copy
- OS clears dirty bit when it brings a page to memory from disk
- Evict dirty pages iff not able to find a clean page

Victim proc's pages

	DB	AB
	0	0
	0	0
	0	0
	1	0
	0	0
	1	0
OS →	1	0
	0	0

Clock algorithm (2)

- Optimisation:
 - We should prefer evicting page that has not changed since it was brought back from disk
 - Such pages can just be deleted without doing a copy
- OS clears dirty bit when it brings a page to memory from disk
- Evict dirty pages iff not able to find a clean page

Victim proc's pages

	DB	AB
	0	0
	0	0
	0	0
	1	0
	0	0
	1	0
	1	0
	0	0

OS →


Clock algorithm (2)

- Optimisation:
 - We should prefer evicting page that has not changed since it was brought back from disk
 - Such pages can just be deleted without doing a copy
- OS clears dirty bit when it brings a page to memory from disk
- Evict dirty pages iff not able to find a clean page


Victim proc's pages

	DB	AB
	0	0
	0	0
	0	0
	1	0
	0	0
	1	0
	1	0
OS →	0	0


Thrashing: Library analogy




Thrashing: Library analogy




Thrashing: Library analogy




Thrashing: Library analogy




Thrashing: Library analogy




Thrashing: Library analogy




Thrashing: Library analogy



Thrashing: Library analogy



Thrashing: Library analogy



Thrashing: Library analogy (2)

- Problem: Library only allows one book to be checked out. Student is constantly running to/from the library. Not able to do work on any assignment.

Thrashing: Library analogy (2)

- Problem: Library only allows one book to be checked out. Student is constantly running to/from the library. Not able to do work on any assignment.
- Solution:

Thrashing: Library analogy (2)

- Problem: Library only allows one book to be checked out. Student is constantly running to/from the library. Not able to do work on any assignment.
- Solution:
 - Reduce “working set”

Thrashing: Library analogy (2)

- Problem: Library only allows one book to be checked out. Student is constantly running to/from the library. Not able to do work on any assignment.
- Solution:
 - Reduce “working set”
 - I will work on OS assignment completely before worrying about DB assignment

Thrashing: Library analogy (2)

- Problem: Library only allows one book to be checked out. Student is constantly running to/from the library. Not able to do work on any assignment.
- Solution:
 - Reduce “working set”
 - I will work on OS assignment completely before worrying about DB assignment
 - I will not work on DB assignment

Thrashing: Library analogy (2)

- Problem: Library only allows one book to be checked out. Student is constantly running to/from the library. Not able to do work on any assignment.
- Solution:
 - Reduce “working set”
 - I will work on OS assignment completely before worrying about DB assignment
 - I will not work on DB assignment
 - Buy a book to avoid going to library

Thrashing

- Total working set of running processes is larger than physical memory. Constantly swapping in and out pages to/from disk. Not able to work.

Thrashing

- Total working set of running processes is larger than physical memory. Constantly swapping in and out pages to/from disk. Not able to work.
- Solution:

Thrashing

- Total working set of running processes is larger than physical memory. Constantly swapping in and out pages to/from disk. Not able to work.
- Solution:
 - Reduce working set

Thrashing

- Total working set of running processes is larger than physical memory. Constantly swapping in and out pages to/from disk. Not able to work.
- Solution:
 - Reduce working set
 - Admission control: run some processes for some time and then some other

Thrashing

- Total working set of running processes is larger than physical memory. Constantly swapping in and out pages to/from disk. Not able to work.
- Solution:
 - Reduce working set
 - Admission control: run some processes for some time and then some other
 - Out-of-memory killer: Linux kills most memory intensive process


Thrashing

- Total working set of running processes is larger than physical memory. Constantly swapping in and out pages to/from disk. Not able to work.
- Solution:
 - Reduce working set
 - Admission control: run some processes for some time and then some other
 - Out-of-memory killer: Linux kills most memory intensive process
 - Buy more memory to avoid copying to/from disk

Reducing memory pressure

Kernel same page merging (KSM)


- OS periodically scans a few pages and computes their hash



Reducing memory pressure

Kernel same page merging (KSM)


- OS periodically scans a few pages and computes their hash



Reducing memory pressure

Kernel same page merging (KSM)


- OS periodically scans a few pages and computes their hash



Reducing memory pressure


Kernel same page merging (KSM)

- OS periodically scans a few pages and computes their hash
- If two page hashes are same, deduplicate
 - Change PTE of one process to point to the common page
- Add duplicate page to free list





Shared pages and demand paging

- Need to update multiple PTEs when pages are swapped out/in



Reverse maps

- When swapping in/out a shared page, all PTEs must be updated.
- rmap: PPN -> list[*PTE]



Paging in action in xv6

xv6 book Ch 2

Process address space in xv6




Figure 2-3. Memory layout of a user process with its initial stack.

Process address space in xv6

- Simplifies maintaining `p->sz`




Figure 2-3. Memory layout of a user process with its initial stack.

Process address space in xv6

- Simplifies maintaining `p->sz`
- Heap grows upwards. `malloc` can ask for more memory via `sbrk`




Figure 2-3. Memory layout of a user process with its initial stack.

Process address space in xv6

- Simplifies maintaining `p->sz`
- Heap grows upwards. `malloc` can ask for more memory via `sbrk`
- Stack is bounded to 4KB




Figure 2-3. Memory layout of a user process with its initial stack.

Guard page

- Present bit is not set for the guard page
- Stack overflow traps into the OS
- OS will kill the process






Figure 3-14. Format of Page-Directory and Page-Table Entries for 4-KByte Pages and 32-Bit Physical Addresses

Paging in xv6



Mapping OS into process address space

- User/supervisor bit is not set for OS pages





Figure 3-14. Format of Page-Directory and Page-Table Entries for 4-KByte Pages and 32-Bit Physical Addresses

Mapping OS into process address space

- User/supervisor bit is not set for OS pages










Figure 3-14. Format of Page-Directory and Page-Table Entries for 4-KByte Pages and 32-Bit Physical Addresses















Advantages of mapping OS into process virtual memory

Trap handling


- Kernel stack and IDT are in the address space of the process. => Hardware need not switch page tables for handling interrupts



Advantages of mapping OS into process virtual memory

Trap handling


- Kernel stack and IDT are in the address space of the process. => Hardware need not switch page tables for handling interrupts



Advantages of mapping OS into process virtual memory

Trap handling


- Kernel stack and IDT are in the address space of the process. => Hardware need not switch page tables for handling interrupts



Advantages of mapping OS into process virtual memory


Read sys call parameters without switching page table

```
int fetchint(uint addr, int *ip) {  
    if(addr >= p->sz || addr+4 > p->sz)  
        return -1;  
    *ip = *(int*)(addr + p->offset);  
}  
  
int argint(int n, int *ip) {  
    return fetchint((myproc()->tf->esp)  
        + 4 + 4*n, ip);  
}
```




Kernel page table isolation (KPTI)

- Need to switch page tables to handle interrupts/syscalls



Kernel page table isolation (KPTI)

- Need to switch page tables to handle interrupts/syscalls
- Syscall and interrupt heavy workloads like Postgres see 7-17% overheads (16-23% without tagged TLB)



Context switching

```
void seginit(void) {  
    c->gdt[SEG_UCODE] = SEG(.., 0, 0xffffffff, DPL_USER);  
    c->gdt[SEG_UDATA] = SEG(.., 0, 0xffffffff, DPL_USER);  
}
```

- User segments map to entire memory since protection is done via paging


Context switching


```
void seginit(void) {  
    c->gdt[SEG_UCODE] = SEG(.., 0, 0xffffffff, DPL_USER);  
    c->gdt[SEG_UDATA] = SEG(.., 0, 0xffffffff, DPL_USER);  
}
```


- User segments map to entire memory since protection is done via paging
- switchvm changes page tables instead of GDT entries

```
void switchvm(struct proc *p) {  
    mycpu()->gdt[SEG_TSS] = SEG16(STS_T32A, &mycpu()->ts,  
                                   sizeof(mycpu()->ts)-1, 0);  
    mycpu()->ts.ss0 = SEG_KDATA << 3;  
    mycpu()->ts.esp0 = (uint)p->kstack + KSTACKSIZE;  
    ltr(SEG_TSS << 3);  
    lcr3(V2P(p->pgdir)); // switch to process address space  
}
```


```
void scheduler(void) {  
    // pick RUNNABLE process p  
    switchvm(p);  
    swtch(p->context);  
}
```





Boot up sequence: BIOS to bootloader to OS



Kernel has different physical and virtual addresses

- kernel.ld declares virtual address `0x80100000`, physical address `0x100000`
- kernel.ld marks `_start` as entry point. `_start` is `V2P_WO(entry)` i.e, `(0x8010000c - 0x80000000)`

```
$ readelf -l kernel
```

Elf file type is EXEC (Executable file)

Entry point `0x10000c`

There are 3 program headers, starting at offset 52

Program Headers:

Type	Offset	VirtAddr	PhysAddr	FileSiz	MemSiz	Flg	Align
LOAD	0x001000	<code>0x80100000</code>	<code>0x00100000</code>	0x07aab	0x07aab	R E	0x1000
LOAD	0x009000	<code>0x80108000</code>	<code>0x00108000</code>	0x02516	0x0d4a8	RW	0x1000
GNU_STACK	0x000000	0x00000000	0x00000000	0x00000	0x00000	RWE	0x10

Section to Segment mapping:

Segment	Sections...
00	.text .rodata
01	.data .bss
02	

```
ENTRY(_start)    kernel.ld
```

```
SECTIONS {
```

```
    . = 0x80100000;
```

```
    .text : AT(0x100000)
```

```
.globl _start    entry.S
```

```
_start = V2P_WO(entry)
```


entry.S sets up an initial page table


eip →

```
entry:
  # Set page directory
  movl    $(V2P_W0(entrypgdir)), %eax
  movl    %eax, %cr3
  # Turn on paging.
  movl    %cr0, %eax
  orl     $(CR0_PG|CR0_WP), %eax
  movl    %eax, %cr0

  movl    $(stack + KSTACKSIZE), %esp
  mov     $main, %eax
  jmp     *%eax

int main (void) {
  kinit1(end, P2V(4*1024*1024));
```

```
__attribute__((__aligned__(PGSIZE)))
pde_t entrypgdir[NPDENTRIES] = {
  // Map VA's [0, 4MB) to PA's [0, 4MB)
  [0] = (0) | PTE_P | PTE_W | PTE_PS,
  // Map VA's [KERNBASE, KERNBASE+4MB) to PA's [0, 4MB)
  [KERNBASE>>PDXSHIFT] = (0) | PTE_P | PTE_W | PTE_PS,
};
```



entry.S sets up an initial page table


eip →

```
entry:
  # Set page directory
  movl    $(V2P_W0(entrypgdir)), %eax
  movl    %eax, %cr3
  # Turn on paging.
  movl    %cr0, %eax
  orl     $(CR0_PG|CR0_WP), %eax
  movl    %eax, %cr0

  movl    $(stack + KSTACKSIZE), %esp
  mov     $main, %eax
  jmp     *%eax

int main (void) {
  kinit1(end, P2V(4*1024*1024));
```

```
__attribute__((__aligned__(PGSIZE)))
pde_t entrypgdir[NPDENTRIES] = {
  // Map VA's [0, 4MB) to PA's [0, 4MB)
  [0] = (0) | PTE_P | PTE_W | PTE_PS,
  // Map VA's [KERNBASE, KERNBASE+4MB) to PA's [0, 4MB)
  [KERNBASE>>PDXSHIFT] = (0) | PTE_P | PTE_W | PTE_PS,
};
```



entry.S sets up an initial page table

eip →


```
entry:
  # Set page directory
  movl    $(V2P_W0(entrypgdir)), %eax
  movl    %eax, %cr3
  # Turn on paging.
  movl    %cr0, %eax
  orl     $(CR0_PG|CR0_WP), %eax
  movl    %eax, %cr0

  movl    $(stack + KSTACKSIZE), %esp
  mov     $main, %eax
  jmp     *%eax

int main (void) {
  kinit1(end, P2V(4*1024*1024));
}
```

CR3 →

```
__attribute__((__aligned__(PGSIZE)))
pde_t entrypgdir[NPDENTRIES] = {
  // Map VA's [0, 4MB) to PA's [0, 4MB)
  [0] = (0) | PTE_P | PTE_W | PTE_PS,
  // Map VA's [KERNBASE, KERNBASE+4MB) to PA's [0, 4MB)
  [KERNBASE>>PDXSHIFT] = (0) | PTE_P | PTE_W | PTE_PS,
};
```




entry.S sets up an initial page table

```
entry:
  # Set page directory
  movl $(V2P_W0(entrypgdir)), %eax
  movl %eax, %cr3
  # Turn on paging.
  movl %cr0, %eax
  orl $(CR0_PG|CR0_WP), %eax
  movl %eax, %cr0

  movl $(stack + KSTACKSIZE), %esp
  mov $main, %eax
  jmp *%eax

int main (void) {
  kinit1(end, P2V(4*1024*1024));
```

```
__attribute__((__aligned__(PGSIZE)))
pde_t entrypgdir[NPDENTRIES] = {
  // Map VA's [0, 4MB) to PA's [0, 4MB)
  [0] = (0) | PTE_P | PTE_W | PTE_PS,
  // Map VA's [KERNBASE, KERNBASE+4MB) to PA's [0, 4MB)
  [KERNBASE>>PDXSHIFT] = (0) | PTE_P | PTE_W | PTE_PS,
};
```




entry.S sets up an initial page table

```
entry:
  # Set page directory
  movl $(V2P_W0(entrypgdir)), %eax
  movl %eax, %cr3
  # Turn on paging.
  movl %cr0, %eax
  orl $(CR0_PG|CR0_WP), %eax
  movl %eax, %cr0

  movl $(stack + KSTACKSIZE), %esp
  mov $main, %eax
  jmp *%eax

int main (void) {
  kinit1(end, P2V(4*1024*1024));
```

```
__attribute__((__aligned__(PGSIZE)))
pde_t entrypgdir[NPDENTRIES] = {
  // Map VA's [0, 4MB) to PA's [0, 4MB)
  [0] = (0) | PTE_P | PTE_W | PTE_PS,
  // Map VA's [KERNBASE, KERNBASE+4MB) to PA's [0, 4MB)
  [KERNBASE>>PDXSHIFT] = (0) | PTE_P | PTE_W | PTE_PS,
};
```




entry.S sets up an initial page table

```
entry:
  # Set page directory
  movl $(V2P_W0(entrypgdir)), %eax
  movl %eax, %cr3
  # Turn on paging.
  movl %cr0, %eax
  orl $(CR0_PG|CR0_WP), %eax
  movl %eax, %cr0

  movl $(stack + KSTACKSIZE), %esp
  mov $main, %eax
  jmp *%eax

int main (void) {
  kinit1(end, P2V(4*1024*1024));
```

```
__attribute__((__aligned__(PGSIZE)))
pde_t entrypgdir[NPDENTRIES] = {
  // Map VA's [0, 4MB) to PA's [0, 4MB)
  [0] = (0) | PTE_P | PTE_W | PTE_PS,
  // Map VA's [KERNBASE, KERNBASE+4MB) to PA's [0, 4MB)
  [KERNBASE>>PDXSHIFT] = (0) | PTE_P | PTE_W | PTE_PS,
};
```




entry.S sets up an initial page table

```
entry:
  # Set page directory
  movl $(V2P_W0(entrypgdir)), %eax
  movl %eax, %cr3
  # Turn on paging.
  movl %cr0, %eax
  orl $(CR0_PG|CR0_WP), %eax
  movl %eax, %cr0


  movl $(stack + KSTACKSIZE), %esp
  mov $main, %eax
  jmp *%eax

int main (void) {
  kinit1(end, P2V(4*1024*1024));
```

```
__attribute__((__aligned__(PGSIZE)))
pde_t entrypgdir[NPDENTRIES] = {
  // Map VA's [0, 4MB) to PA's [0, 4MB)
  [0] = (0) | PTE_P | PTE_W | PTE_PS,
  // Map VA's [KERNBASE, KERNBASE+4MB) to PA's [0, 4MB)
  [KERNBASE>>PDXSHIFT] = (0) | PTE_P | PTE_W | PTE_PS,
};
```



entry.S sets up an initial page table




entry.S sets up an initial page table

```
entry:
    # Set page directory
    movl    $(V2P_W0(entrypgdir)), %eax
    movl    %eax, %cr3
    # Turn on paging.
    movl    %cr0, %eax
    orl     $(CR0_PG|CR0_WP), %eax
    movl    %eax, %cr0

    movl    $(stack + KSTACKSIZE), %esp
    mov     $main, %eax
    jmp     *%eax
```

```
int main (void) {
→ kinit1(end, P2V(4*1024*1024));
```


```
__attribute__((__aligned__(PGSIZE)))
pde_t entrypgdir[NPDENTRIES] = {
    // Map VA's [0, 4MB) to PA's [0, 4MB)
    [0] = (0) | PTE_P | PTE_W | PTE_PS,
    // Map VA's [KERNBASE, KERNBASE+4MB) to PA's [0, 4MB)
    [KERNBASE>>PDXSHIFT] = (0) | PTE_P | PTE_W | PTE_PS,
};
```



main calls kinit1 to initialise free list

Only 4MB is mapped in the initial page table

```
__attribute__((__aligned__(PGSIZE)))
pde_t entrypgdir[NPDENTRIES] = {
    // Map VA's [0, 4MB) to PA's [0, 4MB)
    [0] = (0) | PTE_P | PTE_W | PTE_PS,
    // Map VA's [KERNBASE, KERNBASE+4MB) to PA's [0, 4MB)
    [KERNBASE>>PDXSHIFT] = (0) | PTE_P | PTE_W | PTE_PS,
};
```




main calls kinit1 to initialise free list

Only 4MB is mapped in the initial page table

```
int main (void) {  
    kinit1(end, P2V(4*1024*1024));  
    ...  
}
```

```
__attribute__((__aligned__(PGSIZE)))  
pde_t entrypgdir[NPDENTRIES] = {  
    // Map VA's [0, 4MB) to PA's [0, 4MB)  
    [0] = (0) | PTE_P | PTE_W | PTE_PS,  
    // Map VA's [KERNBASE, KERNBASE+4MB) to PA's [0, 4MB)  
    [KERNBASE>>PDXSHIFT] = (0) | PTE_P | PTE_W | PTE_PS,  
};
```




main calls kinit1 to initialise free list

Only 4MB is mapped in the initial page table

```
int main (void) {  
    kinit1(end, P2V(4*1024*1024));  
    ...  
}
```

```
__attribute__((__aligned__(PGSIZE)))  
pde_t entrypgdir[NPDENTRIES] = {  
    // Map VA's [0, 4MB) to PA's [0, 4MB)  
    [0] = (0) | PTE_P | PTE_W | PTE_PS,  
    // Map VA's [KERNBASE, KERNBASE+4MB) to PA's [0, 4MB)  
    [KERNBASE>>PDXSHIFT] = (0) | PTE_P | PTE_W | PTE_PS,  
};
```




main calls kinit1 to initialise free list

Only 4MB is mapped in the initial page table

```
int main (void) {  
    kinit1(end, P2V(4*1024*1024));  
    ...  
}
```

Now pages can be allocated from the free list!

```
__attribute__((__aligned__(PGSIZE)))  
pde_t entrypgdir[NPDENTRIES] = {  
    // Map VA's [0, 4MB) to PA's [0, 4MB)  
    [0] = (0) | PTE_P | PTE_W | PTE_PS,  
    // Map VA's [KERNBASE, KERNBASE+4MB) to PA's [0, 4MB)  
    [KERNBASE>>PDXSHIFT] = (0) | PTE_P | PTE_W | PTE_PS,  
};
```



OS sets up a new page table

```
int main(void) {  
    kinit1(end, P2V(4*1024*1024));  
    kvmalloc(); // kernel page table  
    ..  
}
```


main.c

Remove identity mapping in low addresses

```
static struct kmap {  
    void *virt;  
    uint phys_start;  
    uint phys_end;  
    int perm;  
} kmap[] = {  
    { (void*)KERNBASE, 0, EXTMEM, PTE_W},  
    { (void*)KERNLINK, V2P(KERNLINK), V2P(data), 0},  
    { (void*)data, V2P(data), PHYSTOP, PTE_W},  
    { (void*)DEVSPACE, DEVSPACE, 0, PTE_W},  
};
```

Mark code as read-only

Map entire physical memory



OS sets up a new page table

```
int main(void) {  
    kinit1(end, P2V(4*1024*1024));  
    kvmalloc(); // kernel page table  
    ..  
}
```


main.c

Remove identity mapping in low addresses


```
static struct kmap {  
    void *virt;  
    uint phys_start;  
    uint phys_end;  
    int perm;  
} kmap[] = {  
    { (void*)KERNBASE, 0, EXTMEM, PTE_W},  
    { (void*)KERNLINK, V2P(KERNLINK), V2P(data), 0},  
    { (void*)data, V2P(data), PHYSTOP, PTE_W},  
    { (void*)DEVSPACE, DEVSPACE, 0, PTE_W},  
};
```

Mark code as read-only

Physical memory



Virtual memory



Map entire physical memory

OS sets up a new page table

```
kmap[] = {  
  { (void*)KERNBASE, 0,          EXTMEM,  PTE_W},  
  { (void*)KERNLINK, V2P(KERNLINK), V2P(data), 0},  
  { (void*)data,      V2P(data),    PHYSTOP, PTE_W},  
  { (void*)DEVSPACE, DEVSPACE,      0,      PTE_W},  
};
```

```
int main(void) {  
    kinit1(end, P2V(4*1024*1024));  
    kvmalloc(); // kernel page table  
    ..  
}
```

```
// Allocate one page table for the machine for vm.c  
// the kernel address space for scheduler processes.  
void kvmalloc(void) {  
    kpgdir = setupkvm();  
    switchkvm();  
}  
  
pde_t* setupkvm(void) {  
    pde_t *pgdir;  
    struct kmap *k;  
    pgdir = (pde_t*)kalloc();  
    memset(pgdir, 0, PGSIZE);  
    for(k = kmap; k < &kmap[NELEM(kmap)]; k++)  
        mappages(pgdir, k->...)   
}  
  
// Switch h/w page table register to the kernel-only  
// page table, for when no process is running.  
void switchkvm(void) {  
    lcr3(V2P(kpgdir)); // switch to the kernel page table  
}
```


OS sets up a new page table

```
kmap[] = {  
  { (void*)KERNBASE, 0,          EXTMEM,  PTE_W},  
  { (void*)KERNLINK, V2P(KERNLINK), V2P(data), 0},  
  { (void*)data,      V2P(data),    PHYSTOP, PTE_W},  
  { (void*)DEVSPACE, DEVSPACE,      0,      PTE_W},  
};
```

Allocate a new
page for page

```
int main(void) {  
    kinit1(end, P2V(4*1024*1024));  
    kvmalloc(); // kernel page table  
    ..  
}
```

```
// Allocate one page table for the machine for vm.c  
// the kernel address space for scheduler processes.  
void kvmalloc(void) {  
    kpgdir = setupkvm();  
    switchkvm();  
}  
  
pde_t* setupkvm(void) {  
    pde_t *pgdir;  
    struct kmap *k;  
    pgdir = (pde_t*)kalloc();  
    memset(pgdir, 0, PGSIZE);  
    for(k = kmap; k < &kmap[NELEM(kmap)]; k++)  
        mappages(pgdir, k->...)  
}  
  
// Switch h/w page table register to the kernel-only  
// page table, for when no process is running.  
void switchkvm(void) {  
    lcr3(V2P(kpgdir)); // switch to the kernel page table  
}
```

OS sets up a new page table

```
kmap[] = {  
  { (void*)KERNBASE, 0,          EXTMEM,  PTE_W},  
  { (void*)KERNLINK, V2P(KERNLINK), V2P(data), 0},  
  { (void*)data,      V2P(data),    PHYSTOP,  PTE_W},  
  { (void*)DEVSPACE, DEVSPACE,      0,        PTE_W},  
};
```

Allocate a new
page for page

Add page table
entries to map 4
areas

```
int main(void) {  
  kinit1(end, P2V(4*1024*1024));  
  kvmalloc(); // kernel page table  
  ..  
}
```

```
// Allocate one page table for the machine for vm.c  
// the kernel address space for scheduler processes.  
void kvmalloc(void) {  
  kpgdir = setupkvm();  
  switchkvm();  
}  
  
pde_t* setupkvm(void) {  
  pde_t *pgdir;  
  struct kmap *k;  
  pgdir = (pde_t*)kalloc();  
  memset(pgdir, 0, PGSIZE);  
  for(k = kmap; k < &kmap[NELEM(kmap)]; k++)  
    mappages(pgdir, k->...)  
}  
  
// Switch h/w page table register to the kernel-only  
// page table, for when no process is running.  
void switchkvm(void) {  
  lcr3(V2P(kpgdir)); // switch to the kernel page table  
}
```

OS sets up a new page table

```
kmap[] = {  
  { (void*)KERNBASE, 0,          EXTMEM,  PTE_W},  
  { (void*)KERNLINK, V2P(KERNLINK), V2P(data), 0},  
  { (void*)data,      V2P(data),    PHYSTOP, PTE_W},  
  { (void*)DEVSPACE, DEVSPACE,      0,      PTE_W},  
};
```

Allocate a new
page for page

Add page table
entries to map 4
areas

Change CR3 to
new page table

```
int main(void) {  
  kinit1(end, P2V(4*1024*1024));  
  kvmalloc(); // kernel page table  
  ..  
}
```

```
// Allocate one page table for the machine for vm.c  
// the kernel address space for scheduler processes.  
void kvmalloc(void) {  
  kpgdir = setupkvm();  
  switchkvm();  
}  
  
pde_t* setupkvm(void) {  
  pde_t *pgdir;  
  struct kmap *k;  
  pgdir = (pde_t*)kalloc();  
  memset(pgdir, 0, PGSIZE);  
  for(k = kmap; k < &kmap[NELEM(kmap)]; k++)  
    mappages(pgdir, k->...)  
}  
  
// Switch h/w page table register to the kernel-only  
// page table, for when no process is running.  
void switchkvm(void) {  
  lcr3(V2P(kpgdir)); // switch to the kernel page table  
}
```



Mapping pages

- mappages takes page directory, virtual address, size, physical address, and permissions
 - It calls walkpgdir with alloc=1 (to allocate page table pages if they do not exist) to find the page table entry
 - It puts physical address in the PTE, marks it as present, and puts other permissions on PTE
- walkpgdir:
 - mimics hardware's page table walk. It takes first 10 bits to index into page directory to find page table page. It takes next 10 bits to index into page table page. It returns page table entry.
 - If page table page does not exist, it allocates a new page and adds it to page directory

main calls kinit2 to expand free list

```
int main (void) {
    kinit1(end, P2V(4*1024*1024));
    kvmalloc();
    ...
    kinit2(P2V(4*1024*1024), P2V(PHYSTOP));
}
```

Now that full physical memory is mapped, rest of the physical memory is made available to allocator




Setting up processes

- kstack, page table pages need not be contiguous

```
void userinit(){
    p = allocproc();
    p->pgdir = setupkvm();
    ...
}


static struct proc* allocproc(void) {
    p->kstack = kalloc();
    sp = p->kstack + KSTACKSIZE;
    ...
}
```



Key changes from paging

```
void userinit(){
    p->pgdir = setupkvm();
    inituvm(p->pgdir, _binary_initcode_start, (int)_binary_initcode_size);
    p->sz = PGSIZE;
    ...
}
```

```
void inituvm(pde_t *pgdir, char *init, uint sz) {
    char *mem;
    mem = kalloc();
    memset(mem, 0, PGSIZE);
    mappages(pgdir, 0, PGSIZE, V2P(mem), PTE_W|PTE_U);
    memmove(mem, init, sz);
}
```




Key changes from paging

```
void userinit(){  
    p->pgdir = setupkvm();  
    inituvm(p->pgdir, _binary_initcode_start, (int)_binary_initcode_size);  
    p->sz = PGSIZE;  
    ...  
}
```

```
void inituvm(pde_t *pgdir, char *init, uint sz) {  
    char *mem;  
    mem = kalloc();  
    memset(mem, 0, PGSIZE);  
    mappages(pgdir, 0, PGSIZE, V2P(mem), PTE_WIPTE_U);  
    memmove(mem, init, sz);  
}
```

Allocate a page,
clear it




Key changes from paging

```
void userinit(){
    p->pgdir = setupkvm();
    inituvm(p->pgdir, _binary_initcode_start, (int)_binary_initcode_size);
    p->sz = PGSIZE;
    ...
}
```

```
void inituvm(pde_t *pgdir, char *init, uint sz) {
    char *mem;
    mem = kalloc();
    memset(mem, 0, PGSIZE);
    mappages(pgdir, 0, PGSIZE, V2P(mem), PTE_WIPTE_U);
    memmove(mem, init, sz);
}
```

Allocate a page,
clear it

Add page to va=0 in
process' page table. Set
user bit



Key changes from paging


```
void userinit(){
    p->pgdir = setupkvm();
    inituvm(p->pgdir, _binary_initcode_start, (int)_binary_initcode_size);
    p->sz = PGSIZE;
    ...
}
```

```
void inituvm(pde_t *pgdir, char *init, uint sz) {
    char *mem;
    mem = kalloc();
    memset(mem, 0, PGSIZE);
    mappages(pgdir, 0, PGSIZE, V2P(mem), PTE_W|PTE_U);
    memmove(mem, init, sz);
}
```

Allocate a page,
clear it

Copy init binary

Add page to va=0 in
process' page table. Set
user bit



sbrk system call

```
int growproc(int n) {
    uint sz;
    struct proc *curproc = myproc();
    sz = curproc->sz;
    if(n > 0){
        sz = allocuvm(curproc->pgdir, sz, sz + n)
    } else if(n < 0){
        sz = deallocuvm(curproc->pgdir, sz, sz + n)
    }
    curproc->sz = sz;
    switchuvm(curproc);
}
```

Reload CR3 to evict all TLB entries
(some of them maybe dealloc'ed)

```
int sys_sbrk(void) {
    int n;
    if(argint(0, &n) < 0)
        return -1;
    growproc(n)
    ..
}
```

Reject if address space will grow into
the OS area

```
int allocuvm(pde_t *pgdir, uint oldsz, uint newsz) {
    if(newsz >= KERNBASE)
        return 0;
    uint a = PGROUNDUP(oldsz);
    for(; a < newsz; a += PGSIZE){
        char* mem = kalloc();
        memset(mem, 0, PGSIZE);
        mappages(pgdir, (char*)a, PGSIZE, V2P(mem), PTE_WIPTE_U)
    }
}
```

Allocate a page,
clear it, add to the
user space

Summary

- Paging provides significant flexibility: we can spray process pages anywhere. Processes can grow/shrink easily since address space need not be contiguous
- It simplifies allocation and removes external fragmentation since pages are of fixed size
- Multi-level page tables help reduce size of the page tables. TLBs help reduce page table walks
- Address translation hardware is much more complex than segmentation
- Demand paging transparently moves pages between memory and disk
- OS maps itself at fixed area in virtual address space of all processes to ease trap handling, sys call parameter reading, and switching page tables